# MITS6005

# Big Data

# Session 5b

# Data Ingestion

# Data Lake – What is it

## Introduction

- Concept of Data lake took off with the advent of Big Data technologies and remains a fluid evolving concept at this time.
- Data Lake is an enterprise level repository of data on commodity hardware, running Big data applications like HADOOP.
- Data originates from multiple applications in the enterprise and is kept available “As Is” in pre-categorized and pre-manipulated state.
- Raw Data is then refined and made available based on the needs of an organization.
- Data lake Implementation projects originate because of the desire to integrate and store massive data sets at a single centralized location to enable cross functional analytics or to lay the basis for building the functional marts, Departmental sandboxes or enterprise warehouse.

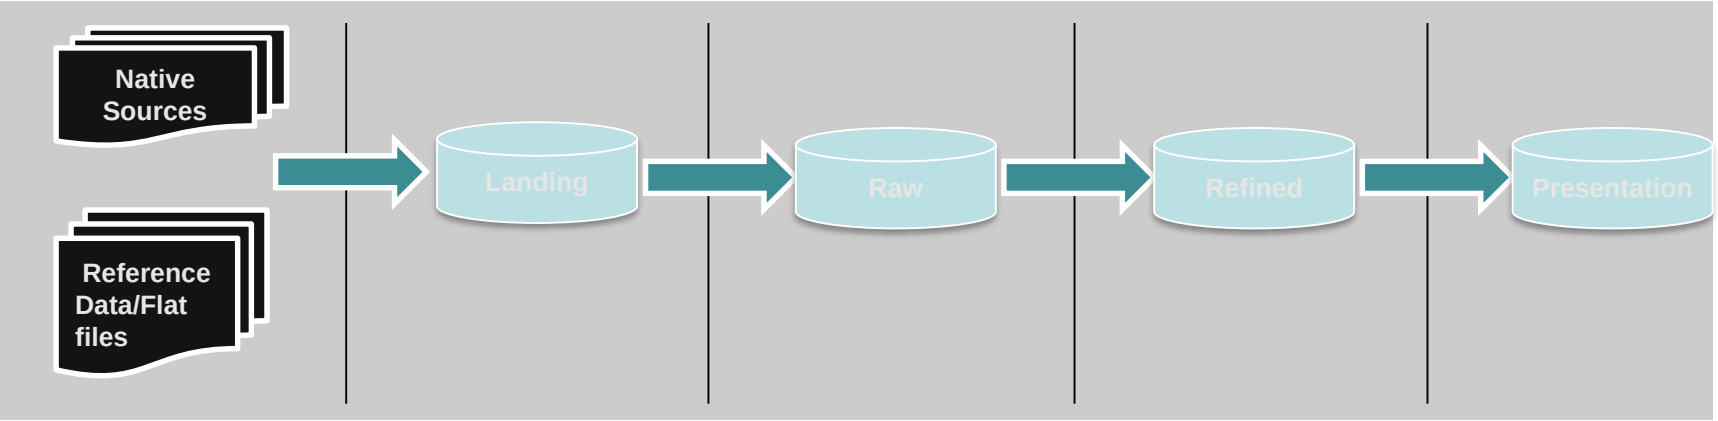

# Value - Unlocking the Value of Data Earlier

Business user community is typically very data "savvy" and will be able to exploit and derive value from data much earlier in the cycle creating new opportunities

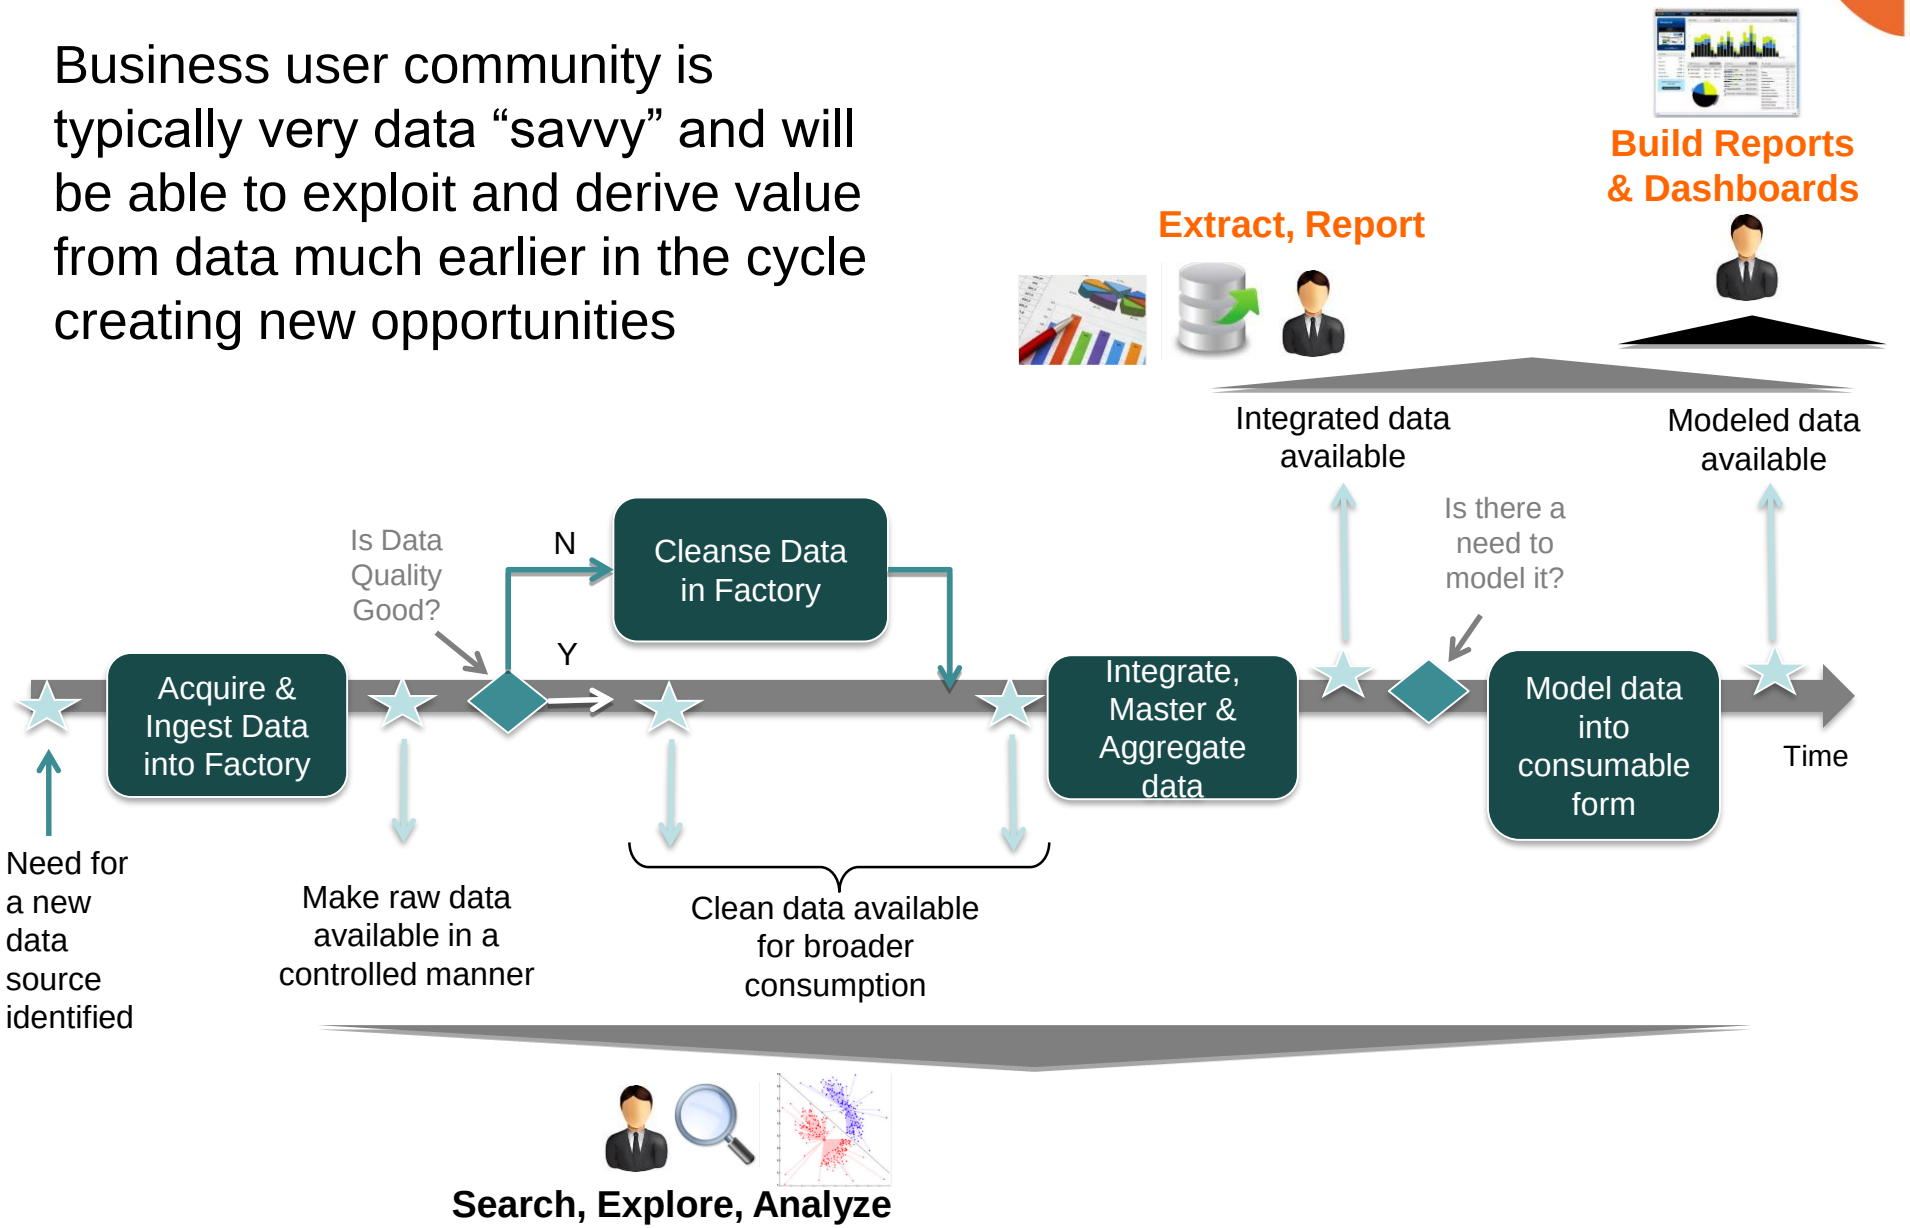

# Data Lake Guiding Principles

- Keep original drivers and objectives “top of mind” and communicate them regularly:
  - “Enable easy integration of new data sources…”
  - “Minimize dependency on costly hardware…”
- Business engagement is essential to understand how data is created and used
- Adapt Incremental implementation approach to Succeed Fast or Fail Fast
- Constantly evaluate tendency to fall back into “old habits”:  Avoid “But that’s how we have always done it…”
- Just-enough data governance necessary to prevent data lakes turning into data swamps
- Select right tool for the job to provide better business value as fast as possible
- Collect metadata for the not only Ingestion automation, but for Data Catalog as a perquisite of Ingesting Data into the Data Lake
- Industry data models must be informed by “the art of the possible” as dictated by source system data structures and business use
- Establish and automate patterns for the ingestion of data into the Data Lake and out of the Data Lake
- Measure data quality upon Ingestion through implementing processes that immediately generate data profiles after ingestion. Create a data quality dashboard in a tool like Tableau to provide visibility into actual data quality

# The Data Lake Paradigm

| | Data Lake | Data Warehouse |
|---|---|---|
| Data Breadth and Depth | • Store Everything As-is, With Complete History<br>• Structured, Semi-Structured and Unstructured | • The Data Warehouse With Aggregated Subsets<br>• Content Management Systems With Limited Metadata |
| Consumption Model | • Let Business Decide What They Need - On-Demand Views<br>• Support Rapid Change | • Pre-Defined Views, Curated By Experts.<br>• Long Change Cycles. |
| Business Driven | • Search Using Business Terminology<br>• Provide Data Lineage and History Tracking and Visualization | • Structured - Tables, Views, Reports. Limited Context<br>• Unstructured - Key-Word Search |
| Data Quality | • Data Quality Is Known And Tracked.<br>• Data Is Available In Various States from Raw to Fully Conformed and Standardized | • Data Available Only After Fully Conformed and Standardized<br>• Quality Metrics Often Not Available |
| Tools | • BYO Data Analysis Tools | • Fixed Set Of Business Intelligence Tools |

# Characteristics of Traditional DA vs. Big Data

| Characteristics | Traditional BI | Big Data |
|---|---|---|
| **Data volume** | Typically Terabytes | Tens to Hundreds of Terabytes, to Petabytes |
| **Velocity of change in scope** | Slower | Faster. Can adapt to frequent change of analytics needs |
| **Total Cost of Ownership** | TOC tends to be expensive | TOC tends to be lower due to lower cost storage and Open source tools |
| **Source data diversity, variety** | Lower | Higher |
| **Analysis driven** | Typically supports known analytics and required reporting | Inherently supports the data analysis and data discovery process by certain users |
| **Requirements driven** | Most of the time | Rarely |
| **Exploration & discovery** | Some of the time | Most of the time |
| **Structure of queries** | Robust | Un-structured |
| **Accuracy of Results** | Deterministic | Approximated |
| **Availability of results** | Slower (longer batch cycles) | Faster |
| **Stored data** | Schema is required to write data | No pre-defined schema is required |

# How things are different

## Data Acquisition Methodologies

### Traditional Data Management


*Data is managed as part of a project, analytic, or other workflow*

---

### "Big Data" Data Management

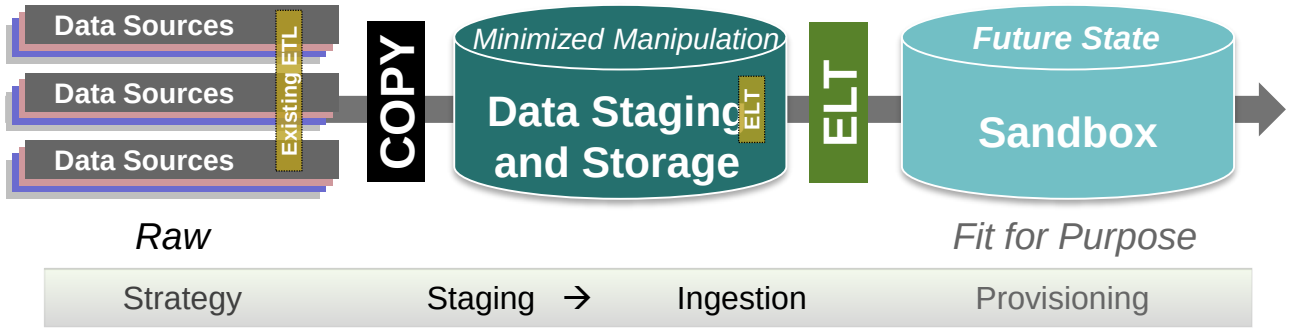


- Data sources are ingested raw (potentially enriched for identifier resolution, searchabilty, quality, etc.)
- Data Services and Access Control move data between storage data stores and consolidate data for analytical data stores

***A principle shared by firms with successful Big Data capabilities is providing as much raw data as possible in an easy to consume, trusted manner***

# Traditional Data Integration

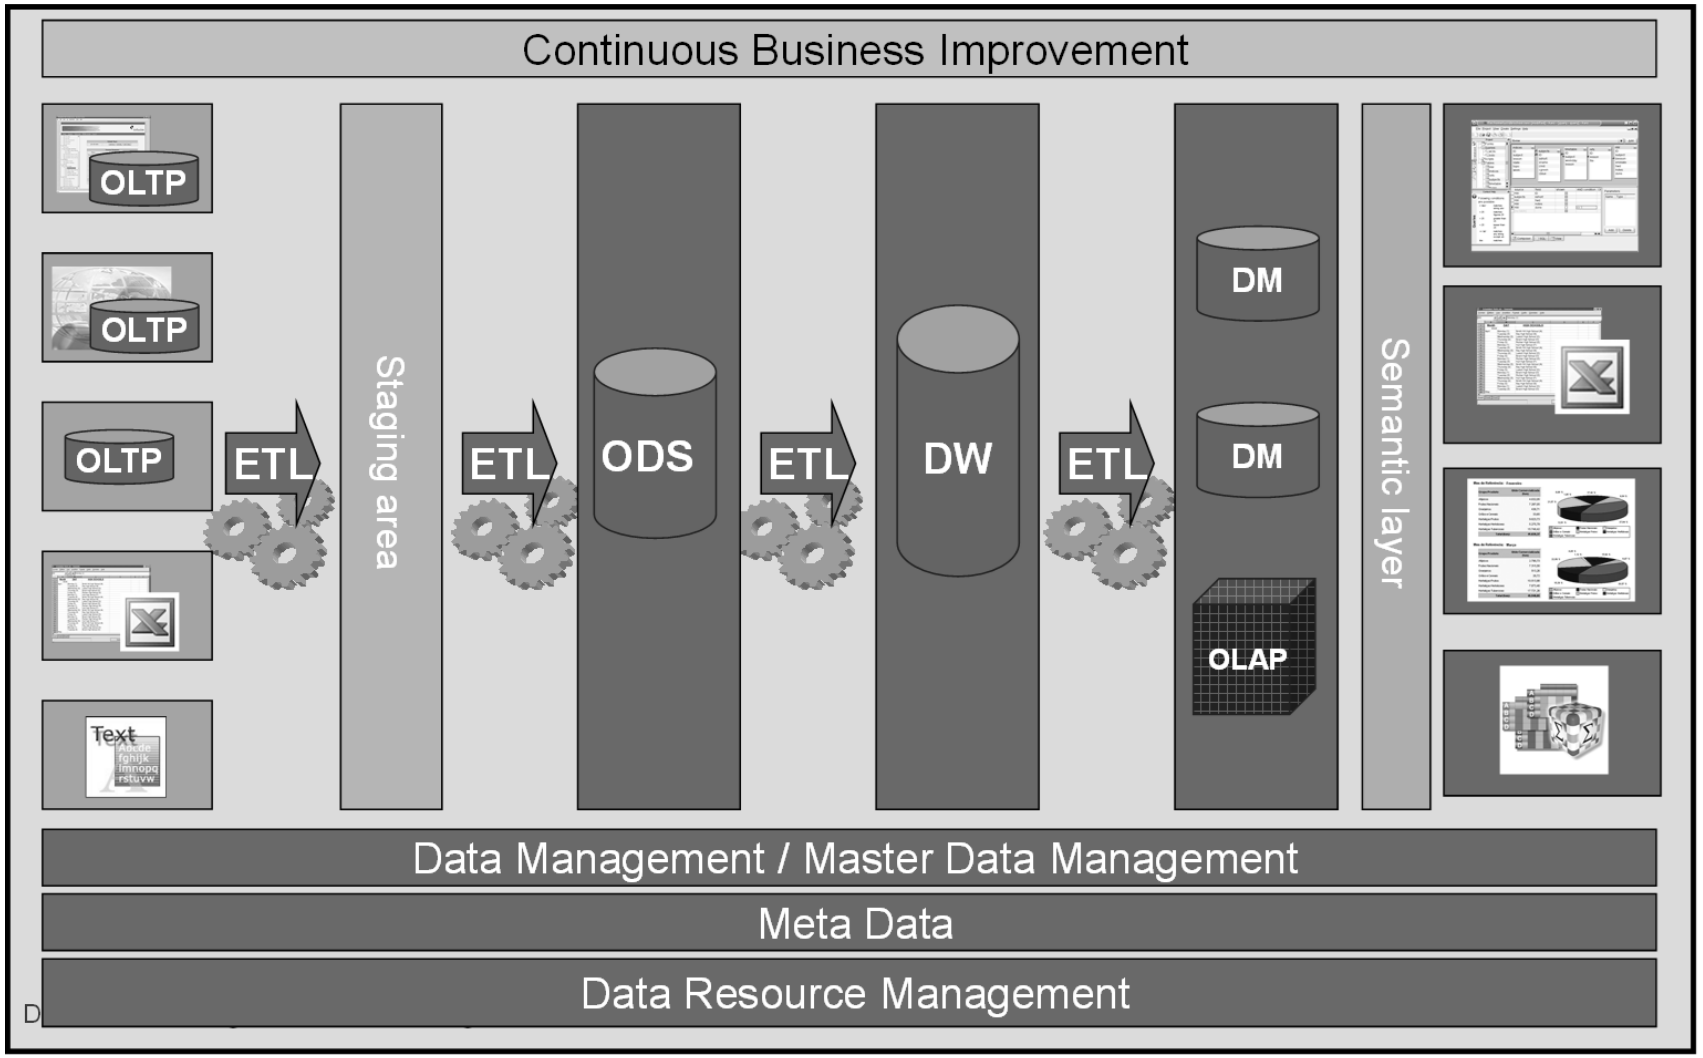

# Traditional Data Integration

**Schema-on-Write (RDBMS):**

- Prescriptive Data Modeling:
  - Create static DB schema
  - Transform data into RDBMS
  - Query data in RDBMS format
- New columns must be added explicitly before new data can propagate into the system.
- Tend to be quite expensive and slow to change
- Limited in terms of the scalability and processing the data as rapidly as the business wants
- **Good for Known Unknowns (Repetition)**

# Modern Day Data Lake Architecture

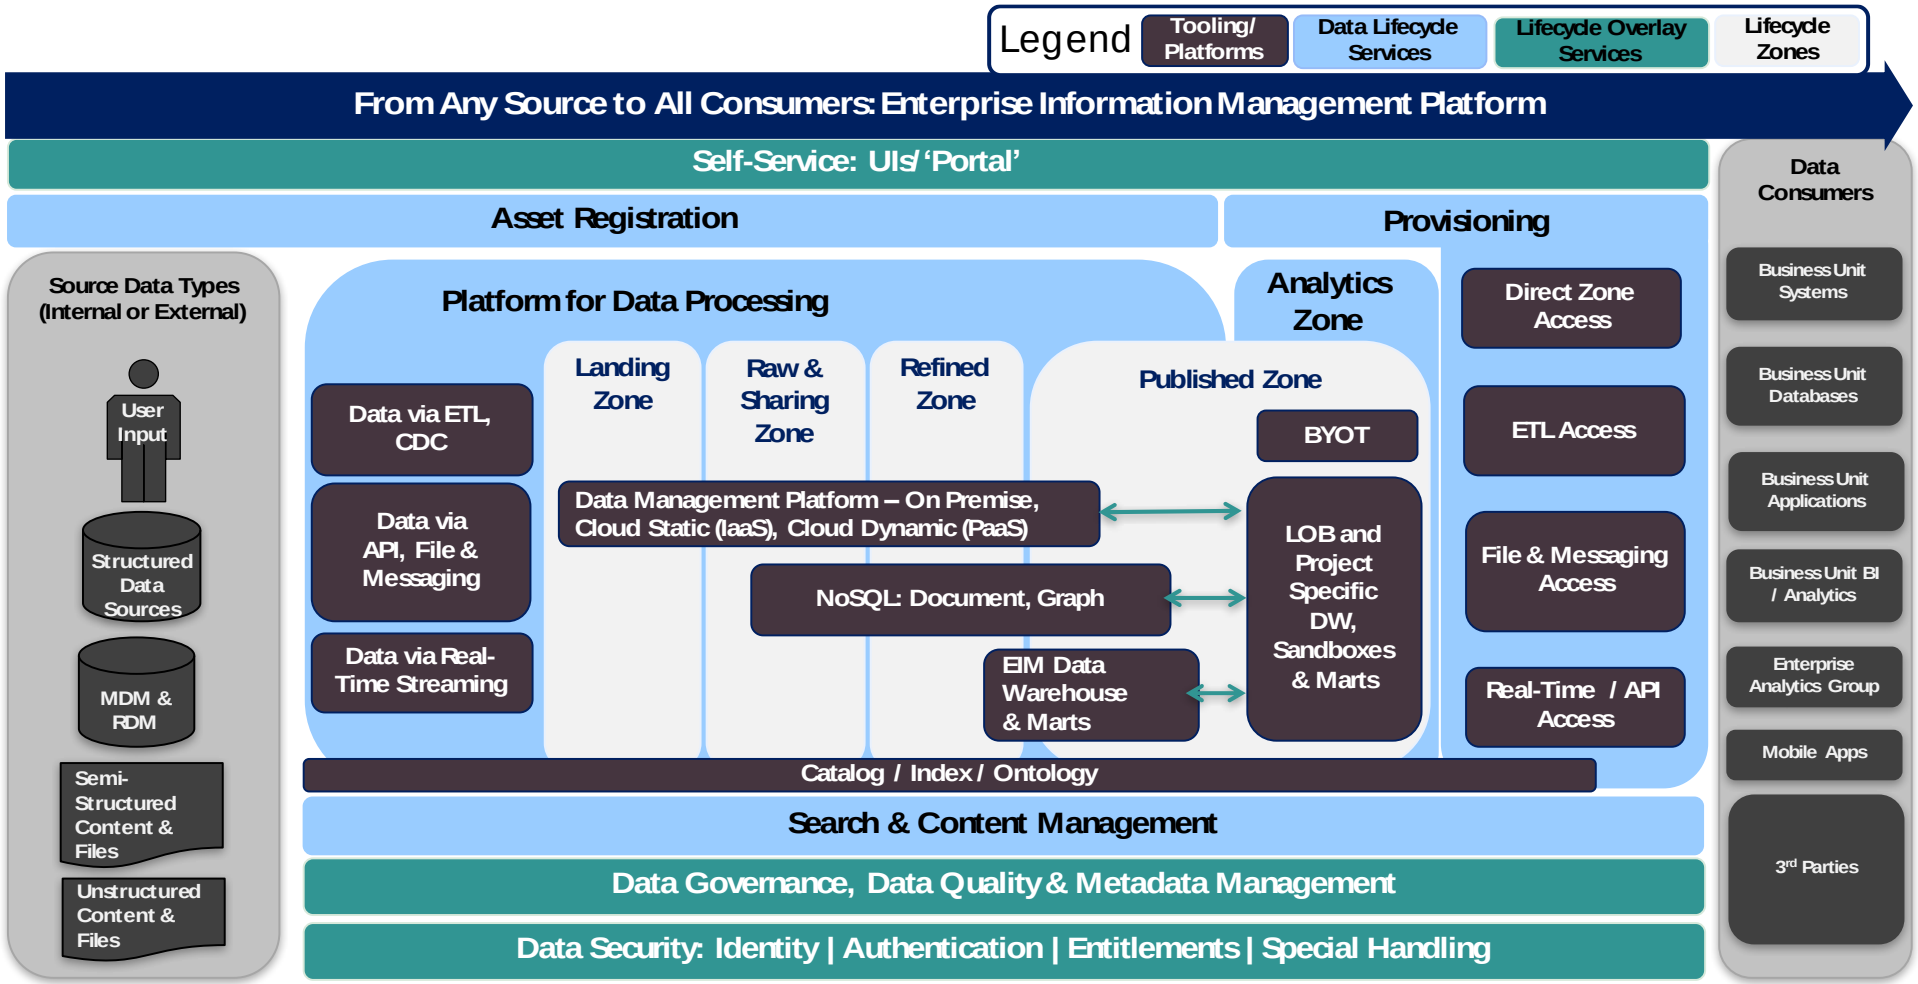

# Modern Day Data Lake Architecture

**Schema-on-Read (Hadoop):**

- Descriptive Data Modeling:
  - Copy data in its native format
  - Create schema + parser
  - Query Data in its native format (does ETL on the fly)

- New data can start flowing any time and will appear retroactively once the schema/parser properly describes it.
- Flexibility and Scalability
- Rapid Data Ingestion
- **Good for Unknown Unknowns (Exploration)**

# Data Lake Integration - Strategy & Planning

Different views or perspectives on the Data Management Architecture will facilitate understanding of recommendations and implications

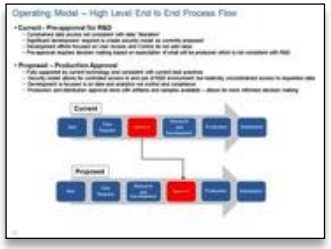

## Business Architecture

- Business Capability View: Business and management processes, their strategic objectives and required data analytics capabilities
- Operating Model / Functional View: Description of key policies, procedures and governance models (committees, review and approval points) required to achieve business objectives
- Organizational View: Description of teams, staffing requirements and reporting relationships in order to support the operational model

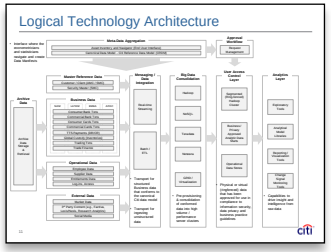


## Technology Architecture

- Logical and physical blueprints for enabling technology capabilities (i.e., Hadoop and/or EDW repositories, Information Asset Inventory & Navigator)
- Vendor strategy and technology product selections

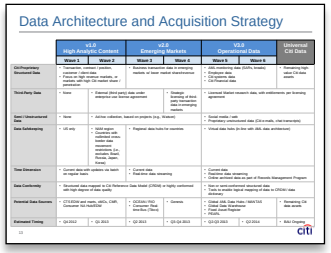


## Data Architecture and Acquisition Strategy

- Identification of structured and unstructured data to expose for analytics, and their data sources
- Strategy / approach for conforming to data dictionary / ontology
- Prioritization and schedule for pre-provisioning data into production environment

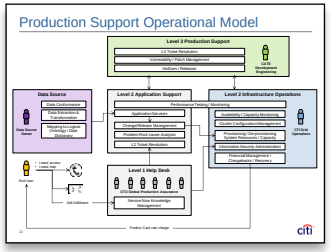


## Physical Infrastructure Support Architecture

- Organizational and process model for how to support end users in a production environment
- Includes model for help desk and ticket resolution, environment monitoring, provisioning and access control management, and financial chargeback/recovery
- Organizational model and Production Service Definition

# Key Capabilities

**Rapid Data Provisioning:** Automated, Metadata configure, data ingestion process sources new data in days

**Data Curation and Quality:** Data quality is measured, tracked, and improved.

**Data Catalog:** Business-focused data dictionary with Google-like search capabilities

**Scalable Data Governance:** Data governance enabled by the architecture.

**Integrated Data Security:** Leverage Tokenization and Encryption to secure critical data

**EDH Platform:** Scalable and Agile Standards-based data management and analytic processing capability

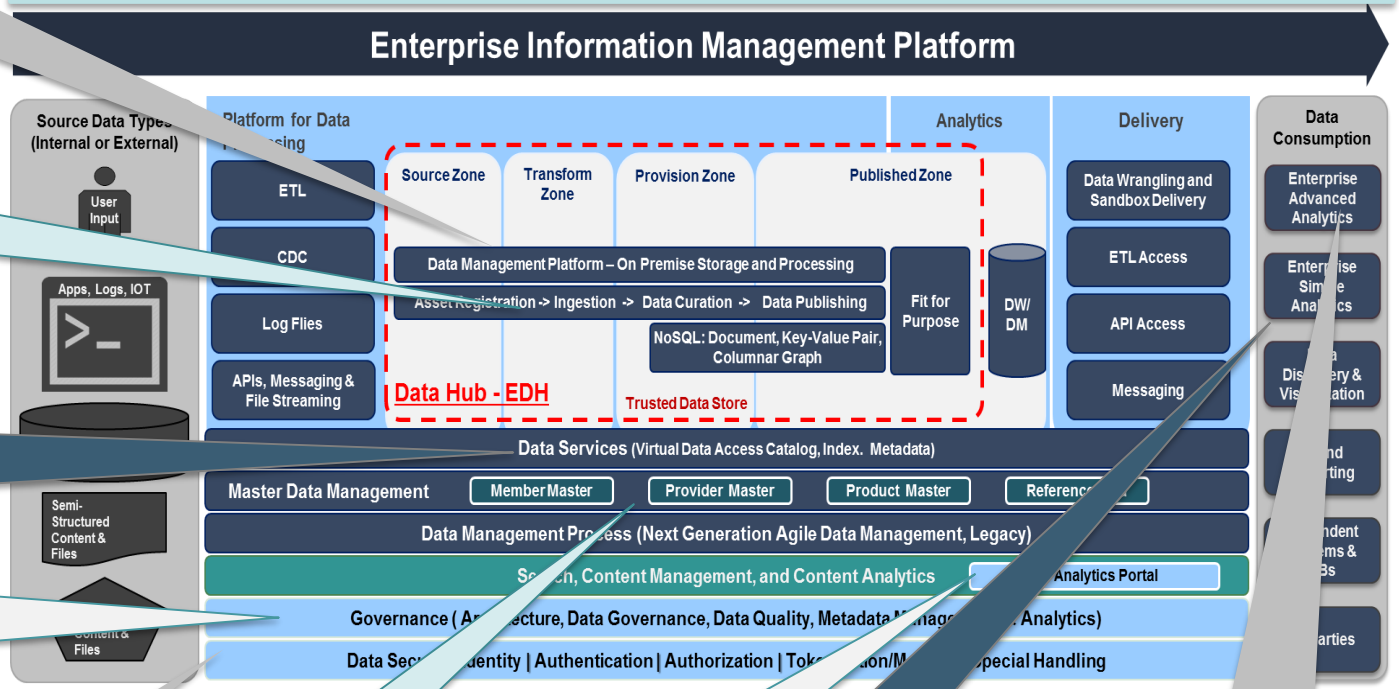


**Holistic View of Member, Provider:** Single view of high-value data.

**Collaborative Knowledge Sharing:** BI/Analytics Portal provides shared access to critical analytical content and best practice information

**Business Authored Reporting and Dashboards:** Data Discovery and Visualization tools enable better understanding of data, thus enabling better corporate performance

**Business Enabled Analytics:** Data Scientists are able to prepare data sets, perform advanced analytics, and publish their results

# Integration - Analytics Driven Approach

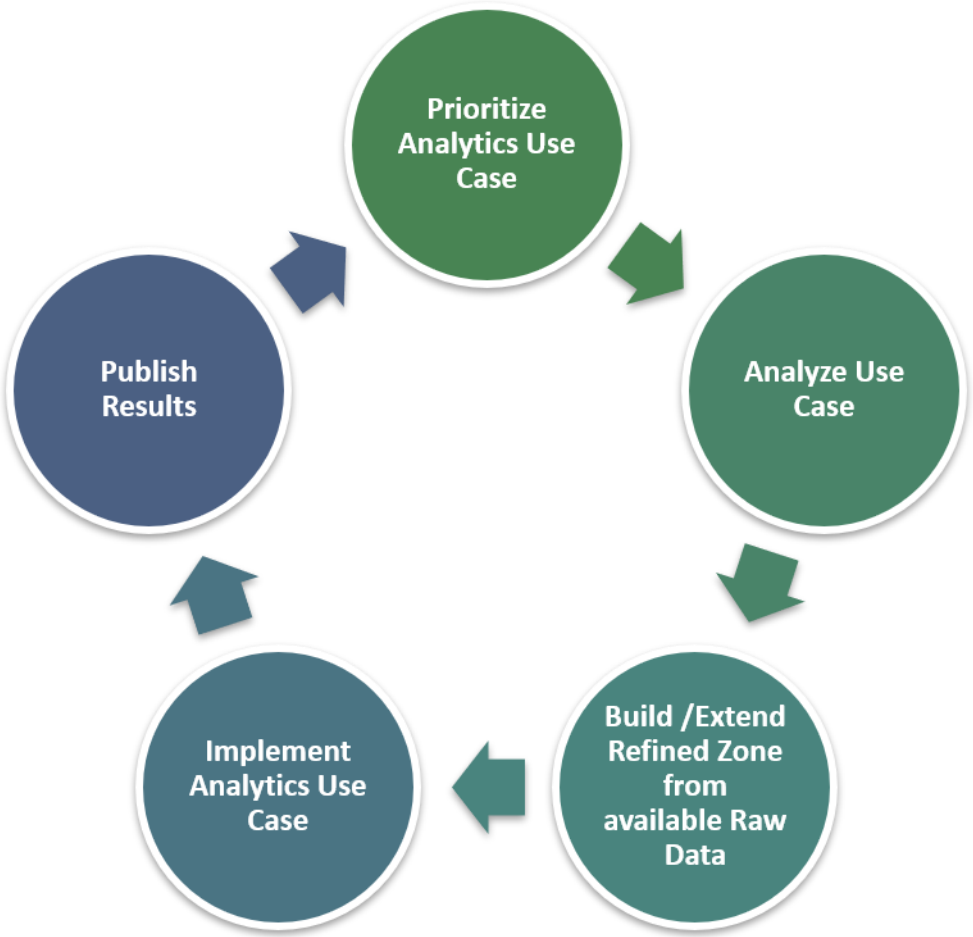


## Key Facts

- Business initiated programs
- Business outcome and value drives priorities
- Analytics-Ready dataset, the priority
- Earlier challenges include cross-domain and cross systems integration
  - Claims matching with Membership
  - Merging Membership from multiple source systems

# Integration - Domain Driven Approach

**Key Facts**

- IT initiated programs
- IT value drives priorities
- Building an Enterprise level data model, a first priority
- Earlier challenges include cross systems integration
  - Merging Membership from multiple source systems

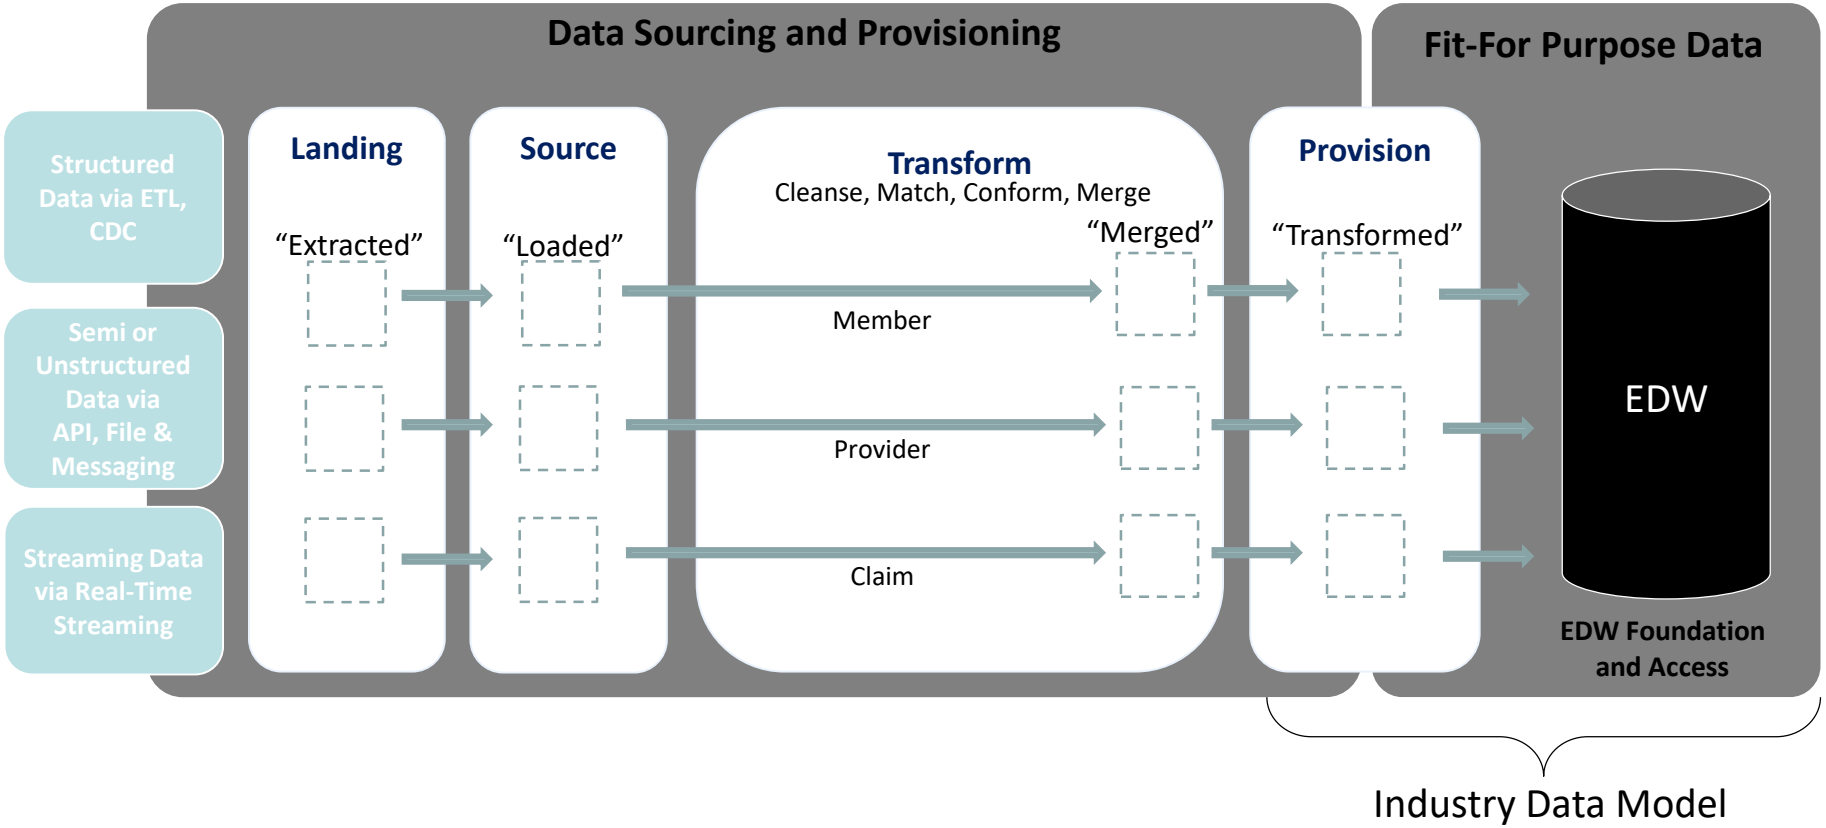

# Agile – Continuous Delivery Approach

**Initiate**

- Requirements Analysis
- Epics/Features Defined Groomed Backlog
- Architecture Defined
- Platform Stood Up
- Testing Approach Defined

**Release Definition**

- 8 Weeks Release
- 5 Two-Week Sprints
- 4 Build Sprints
- 1 HIP Sprint
- HIP Sprint Reserved for
  - Technical Debt
  - Utility Building
  - Sprint Planning

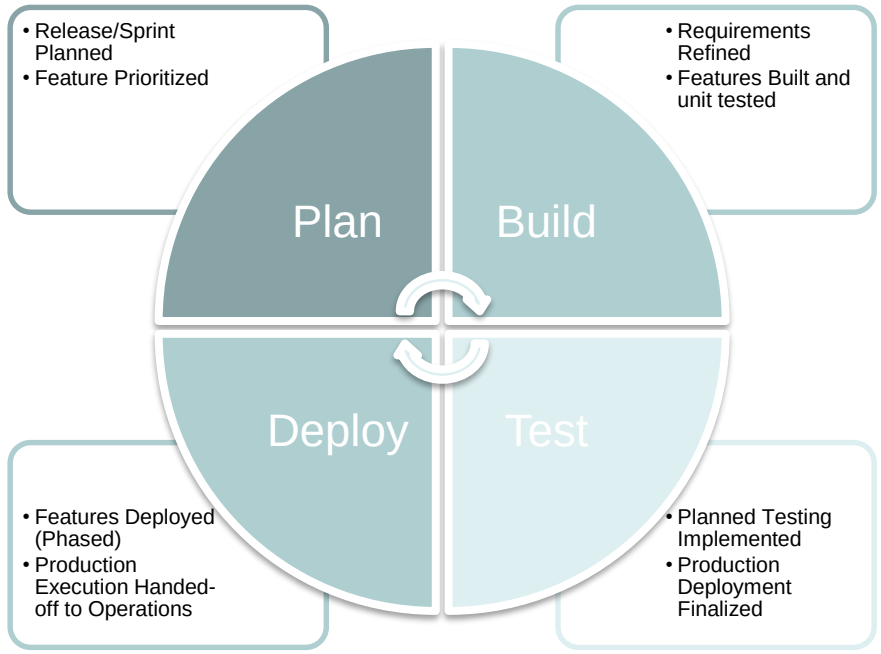


**Typical Release**

Succeed Fast or Fail Fast

# Data Movement & Zones

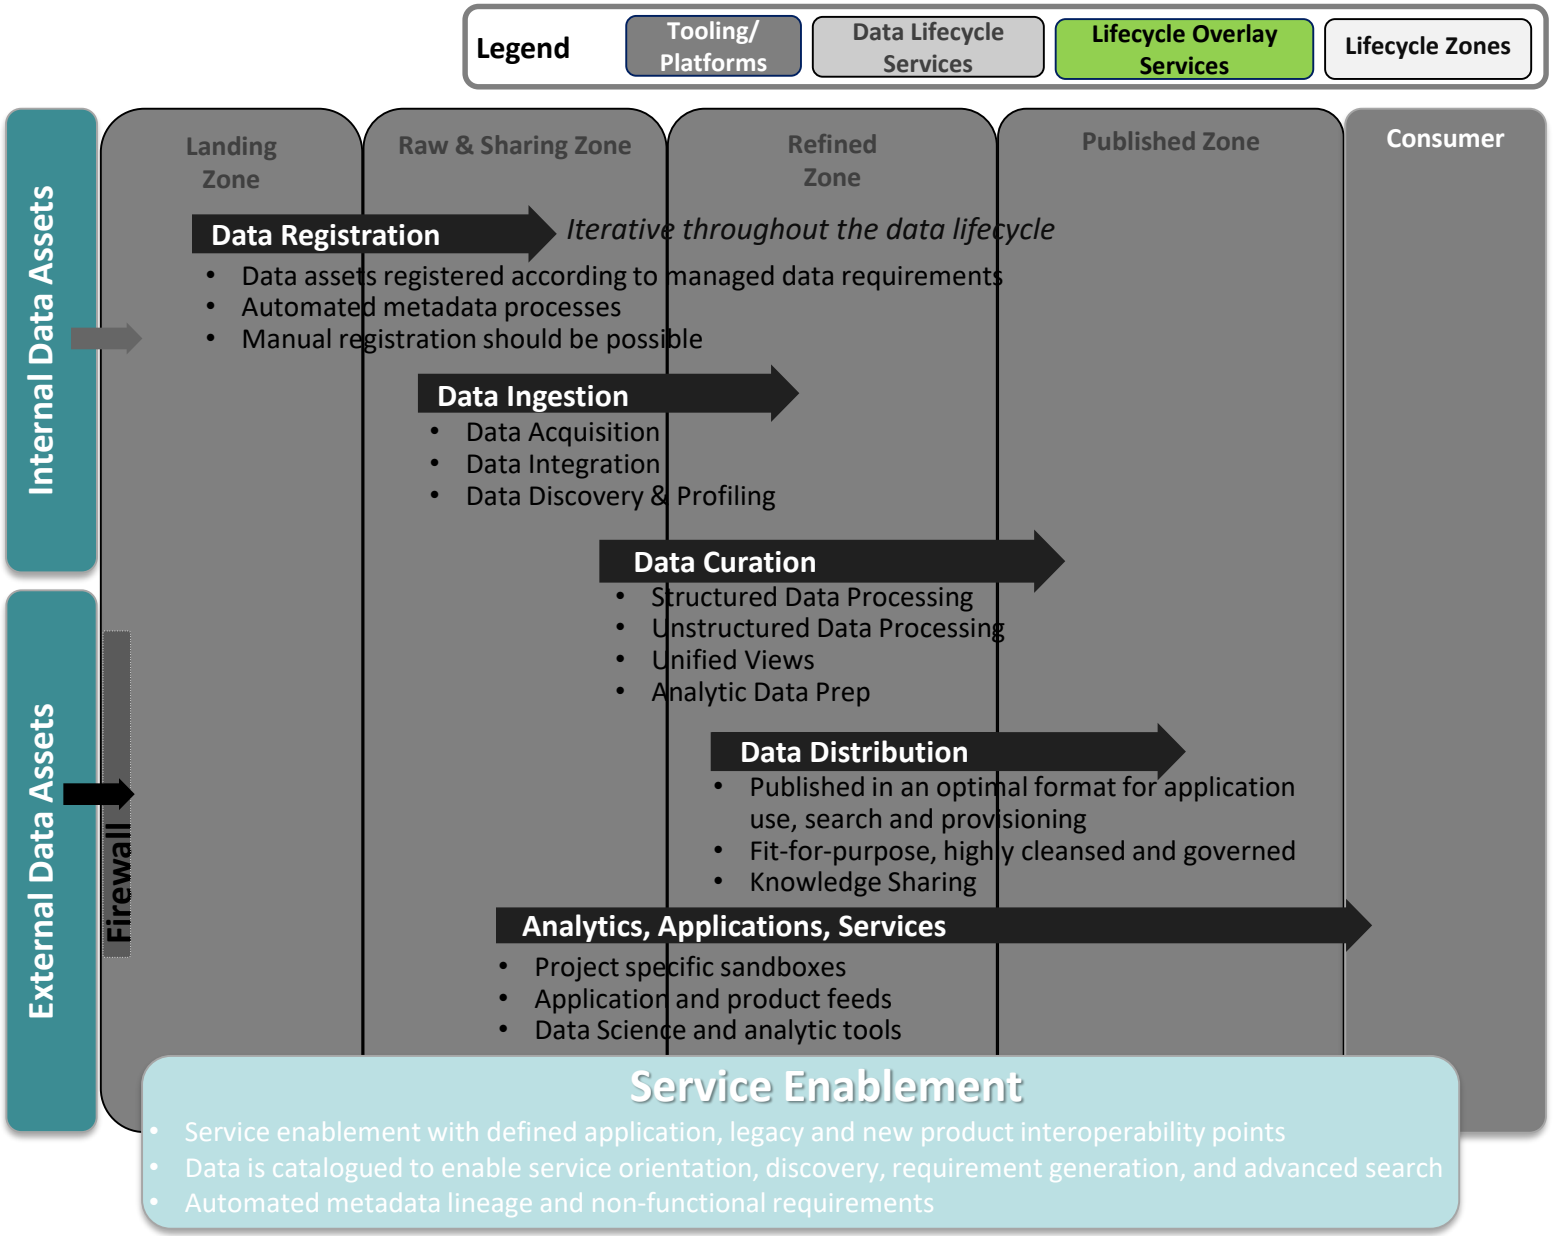

# Data Ingestion

- Process of Moving Data to the Data Lake
- Once Ingested, Data is Available
  - For Processing and Distribution
  - Discovery
  - Analytics
- Key Decisions
  - Streaming vs. Batch
  - Traditional vs. Metadata Driven

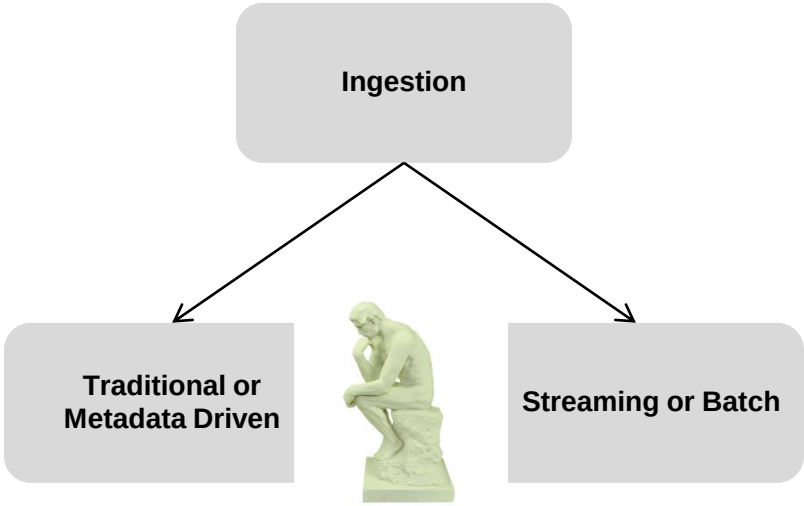

# Batch vs. Streaming

***Batch Processing***

- Similar to Traditional data integration processing
- Good for groups or snapshots of data
- Time lag in data availability

***Stream Processing***

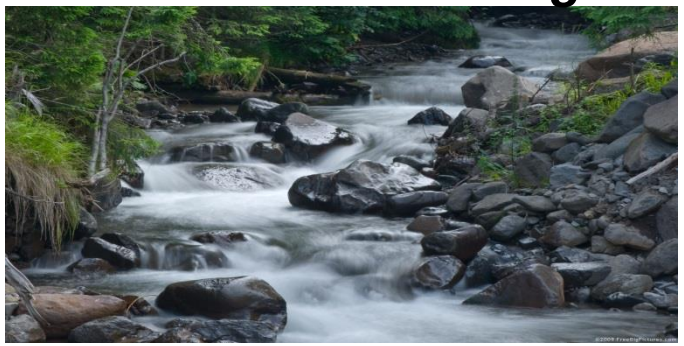

- Fast-data
- Real-time needs
  - Fraud detection
  - Order processing
- Minimal lag in data availability
- Can be more scalable

# Traditional vs. Metadata Driven Ingestion

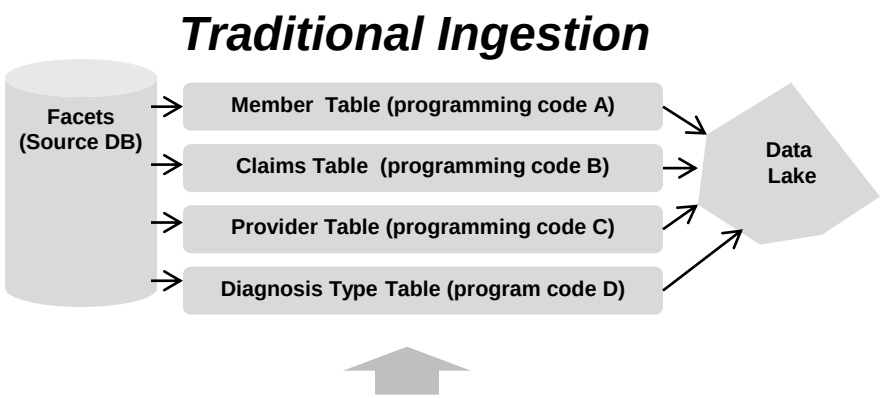


- One set of code for each table ingested
- Good for custom coding or small number of ETL Jobs

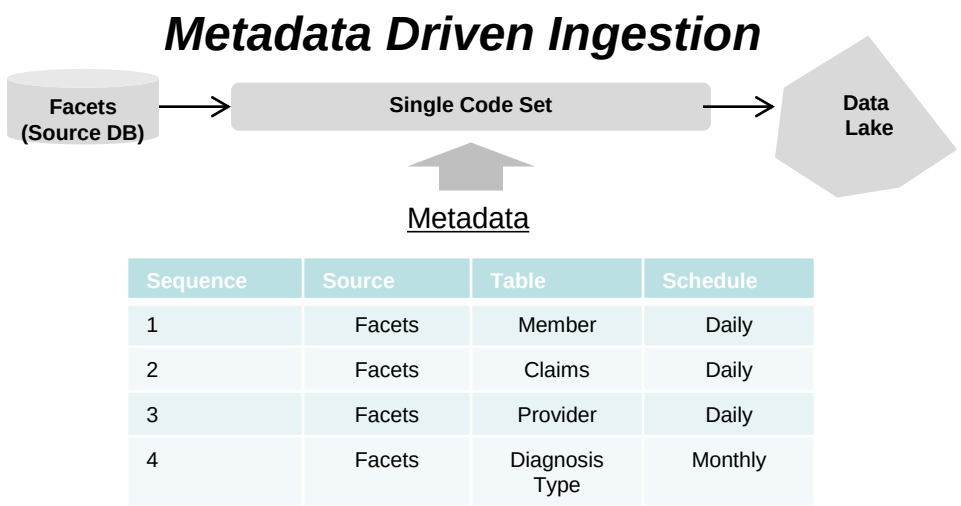


| Sequence | Source | Table | Schedule |
|---|---|---|---|
| 1 | Facets | Member | Daily |
| 2 | Facets | Claims | Daily |
| 3 | Facets | Provider | Daily |
| 4 | Facets | Diagnosis Type | Monthly |

## *Benefits of Metadata Driven Ingestion*

- Extremely scalable (in development time) over 100s and 1000s of tables
- Increased consistency and supportability
- Increased quality of data
- Quick time to value: 5-10x faster  than custom coding or point-to-point usage of ETL tools

# Data Acquisition and Ingestion Detail

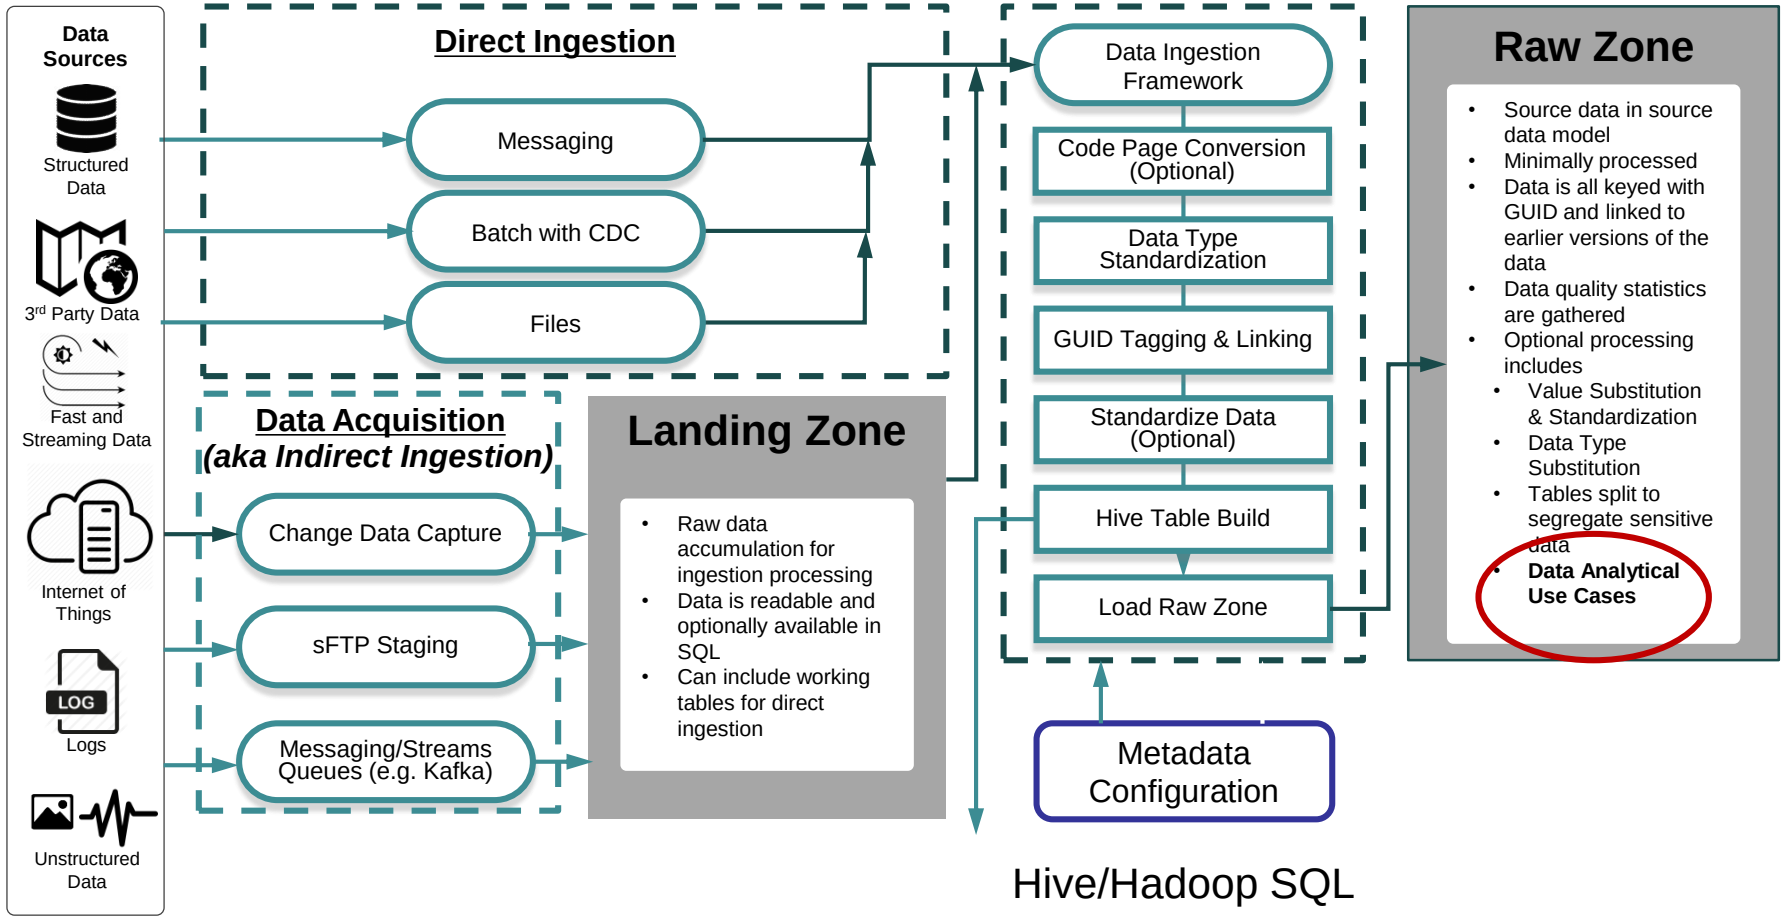

# High-Performance Streaming Ingestion

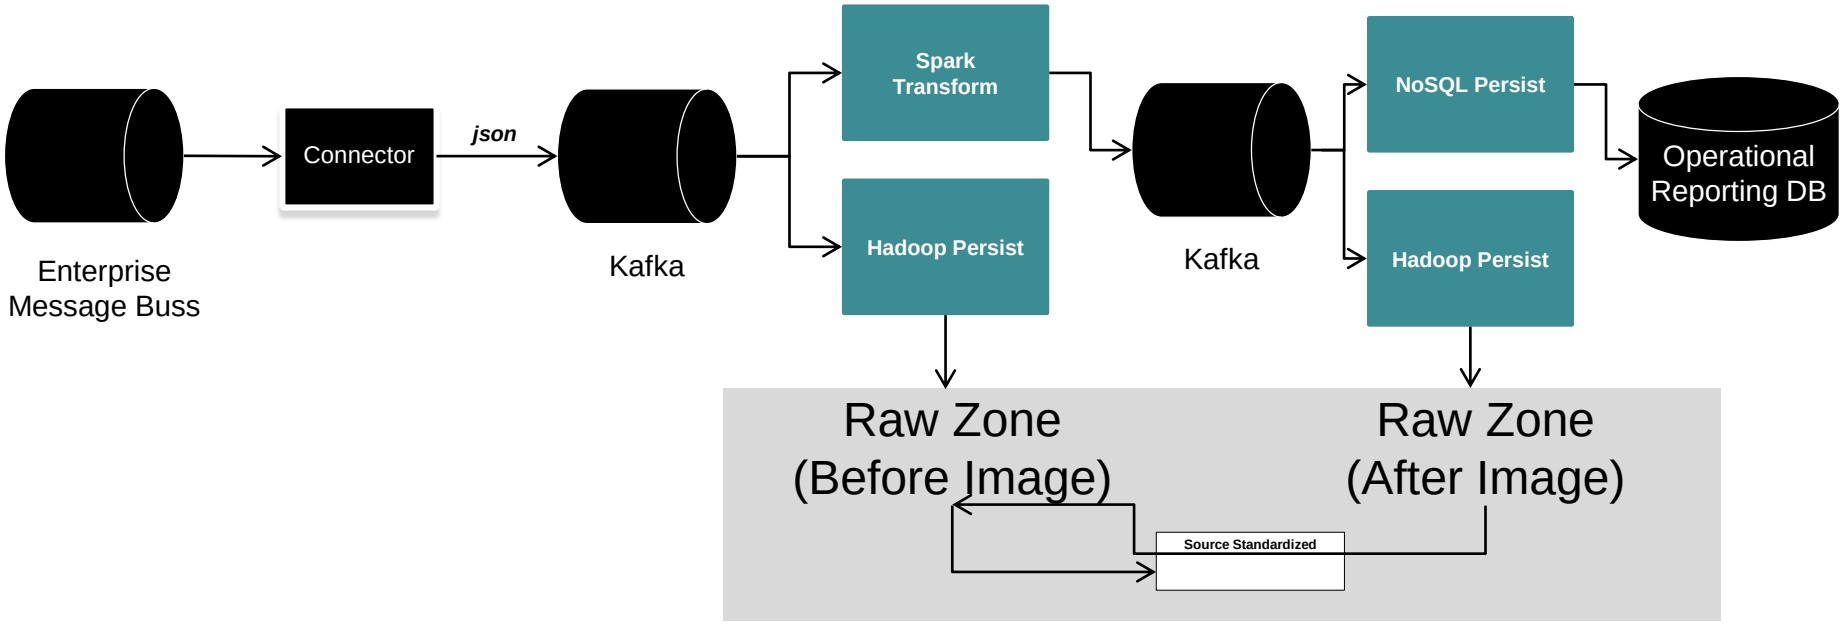

# Catalog - Search & Explore Data

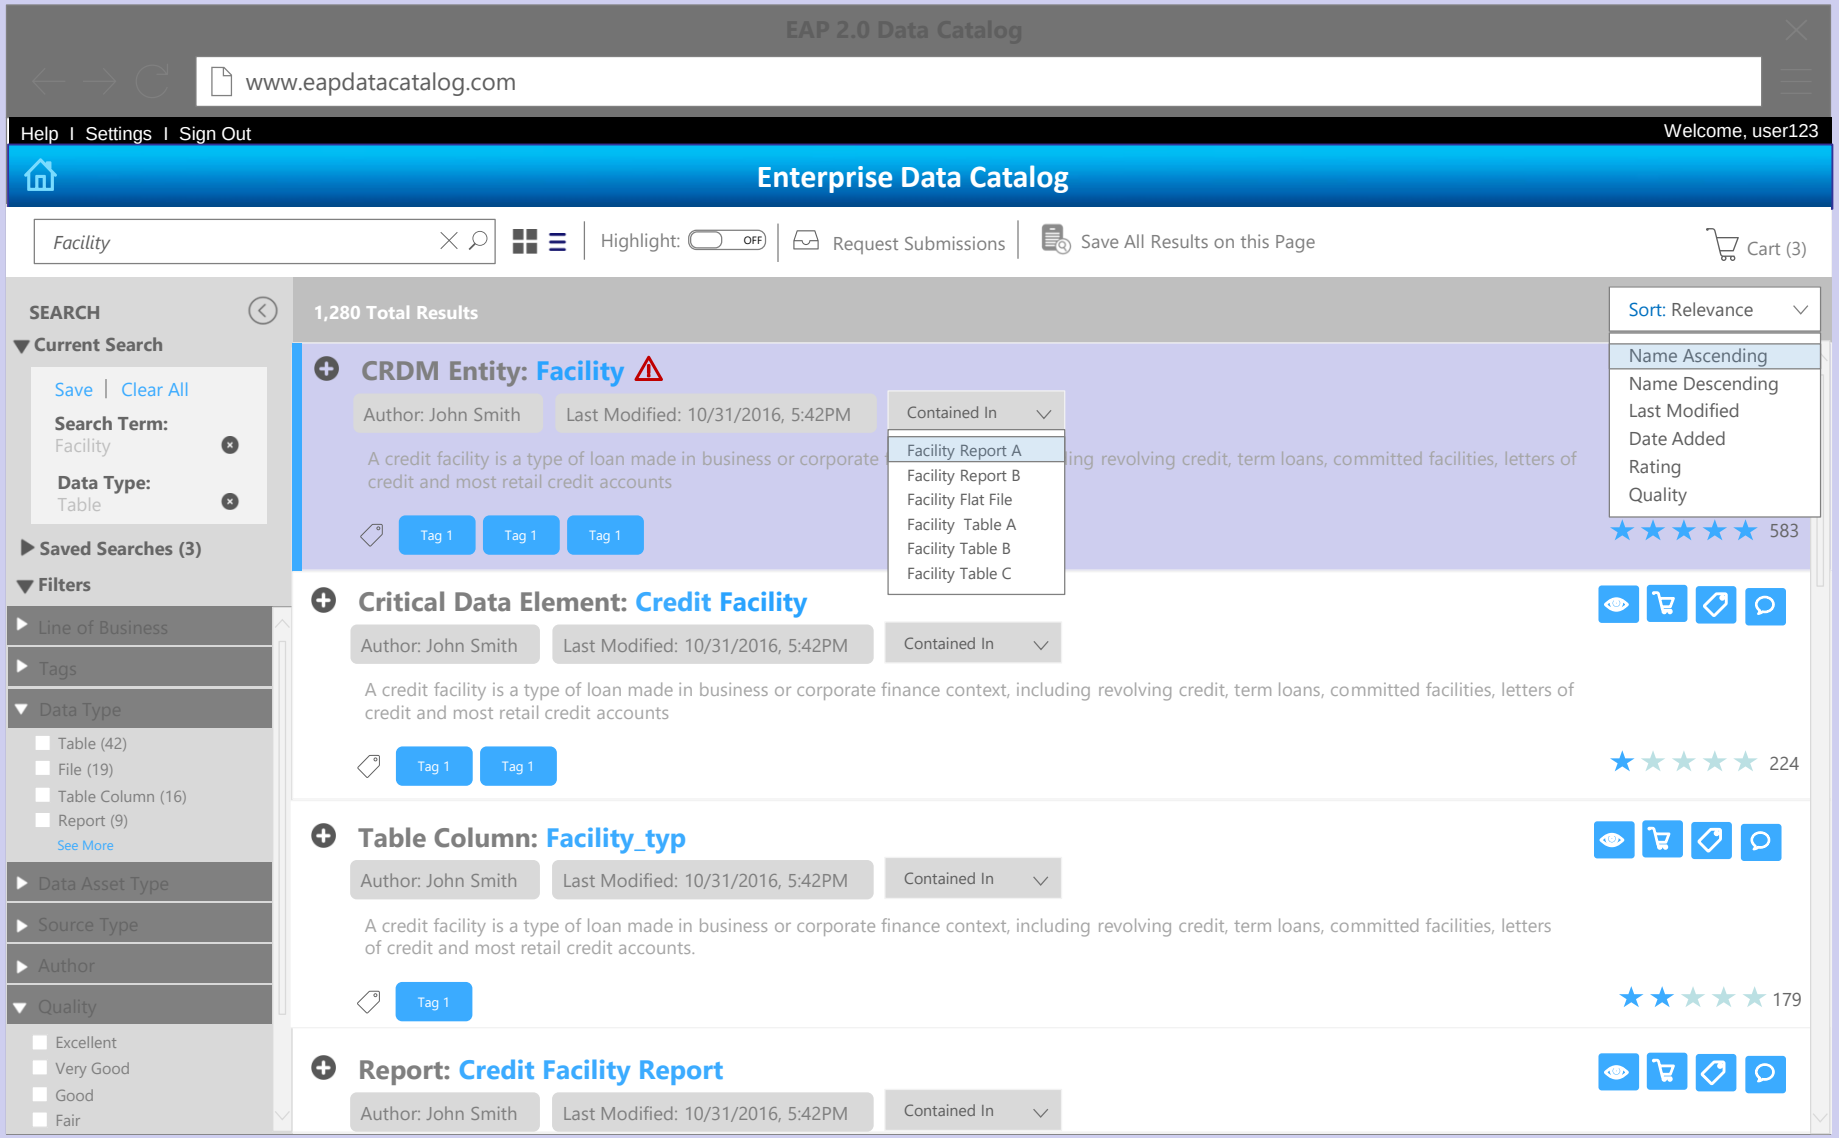

# Masking & Encryption

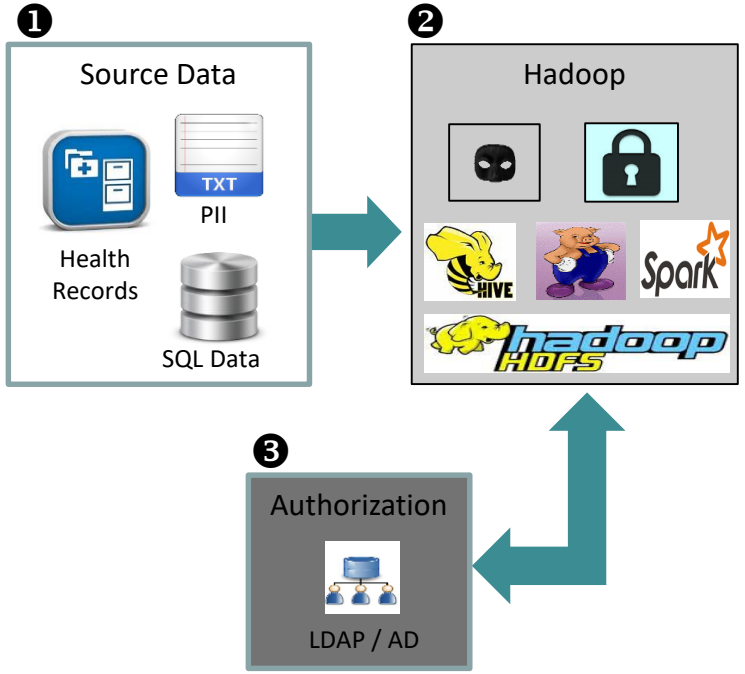


1. Source data in database, data warehouse and various other structures

2. Healthcare firm requires both masking (one-way) and encryption/ decryption for various sensitive data elements

3. Leveraging Active Directory and LDAP, organization controls which users can see what degree of sensitive data– all 100% transparent and automatic to users

# Security & Audit

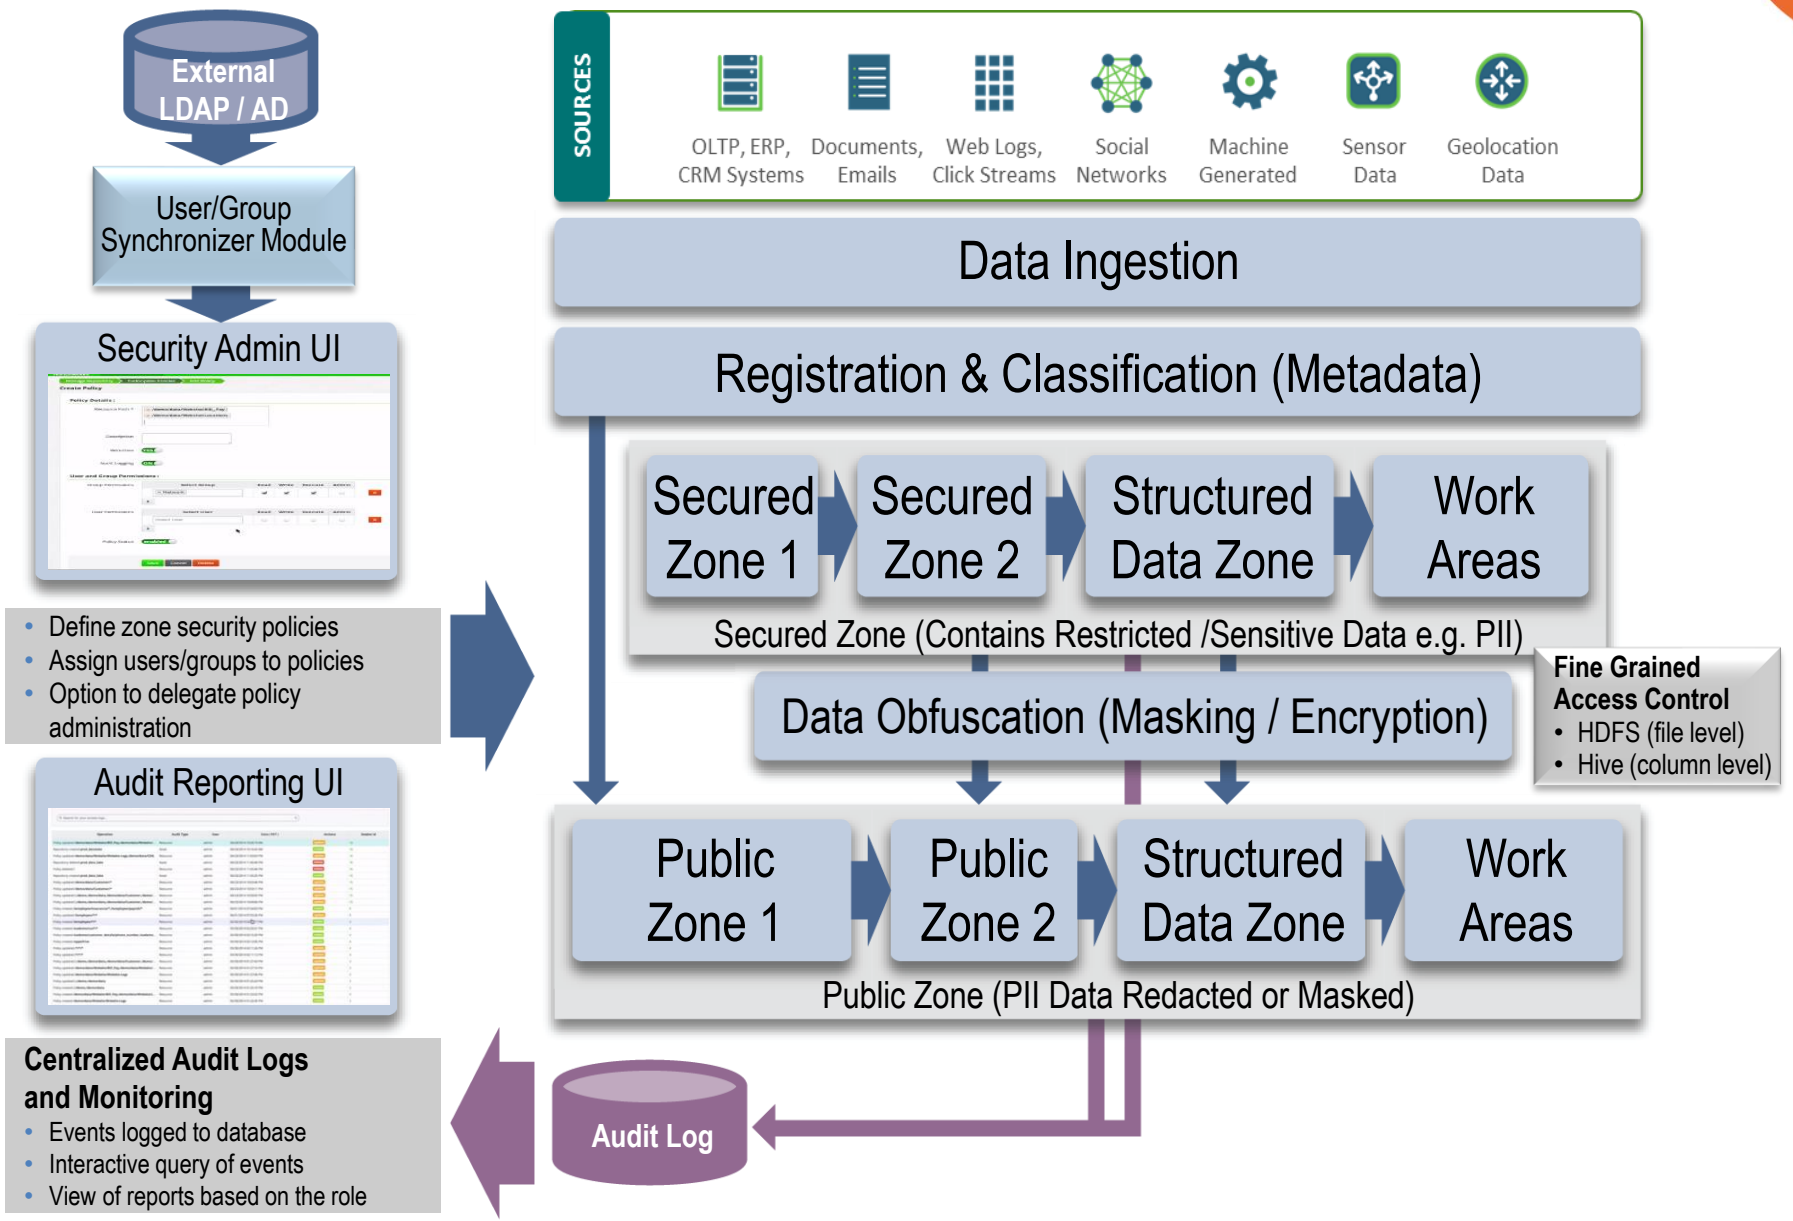

# Users & Governance

**The Zones within the architecture service different user groups within the analytical community**

**Analytical Users**

**Data**

| Zones | Landing | Raw | Refined | Publish |
|---|---|---|---|---|
| **Users within Zones:**<br>• **Actor**<br>• **Tools** | Users:<br>• IT | Users:<br>• Data Scientists<br>• IT<br>• Data stewards<br>Tools: R, Python, SAS | Users:<br>• Data Scientists<br>• Business Power Users<br>Tools: R, Python, SAS | Users:<br>• Business Users<br>• Operational Systems<br>Tools: Qlik/Tableau, Reporting, Service Bus |
| **Data (within zones)**<br>• **Structure** | • Data remains in "as is" form without change<br>• Transient storage; data is removed after ingestion | • As identical as possible to the source data<br>• Data profiled & quality assessed<br>• Metadata augmented with data-specific information derived from discovery, profiling and quality checks<br>• Metadata enriched with business rules & context | • Data cleansed, aggregated, conformed, curated, remediated, or otherwise manipulated according to defined processes and rules<br>• Enriched data driven by business outcomes or analytic needs | • Data Certified for:<br>• Meet specific, defined and managed enterprise data management requirements<br>• High Quality and trusted<br>• Fit For Purpose/Operational Use<br>• Contextually relevant and accurate<br>• Governed by Business specific usage |
| **Data Management (between zones)**<br>• **Governance**<br>• **Security** | Security: Targeted user base, data access rules | Governance:<br>• Data Catalog<br>• Data is Profiled<br>• Quality is Measured<br>Security: Targeted user base, data access rules | Governance:<br>• Lineage Captured<br>• Enterprise Standards<br>• Data Quality Rules<br>• Enterprise Models<br>Security: Targeted user base, data access rules | Governance:<br>• Lineage Captured<br>• Modeled for a Specific Business Use<br>Security: Broad user base, access aligned to limited datasets |

**Governance and security is applied during data movement across zones and within the zones**

# Operating Models

An essential component of enterprise data strategy will be a detailed approach for supporting and operating the future state

**MISSION STATEMENT**


Mission statements and guiding principles for organizational change

**SERVICE MODEL**


Identification and definition of all services to be offered

**ORGANIZATIONAL MODEL**

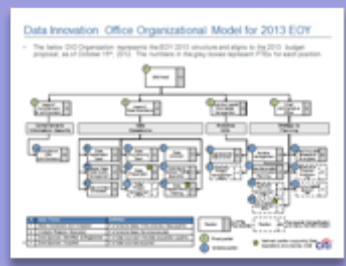


Organizational patterns for Management & Governance, Data, Analytics and Technology

**KEY ARTIFACT TEMPLATES**

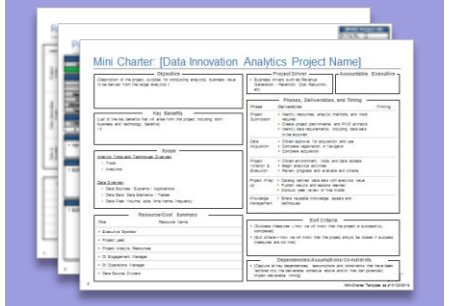


Standard templates to define new projects and track progress

**ROLES & RESPONSIBILITIES**

Roles & Responsibilities for industry standard job descriptions and skill requirements

**PLAYBOOK**


Playbooks that describe workflows for the services and procedures for production support activities

**BUDGET MODEL**

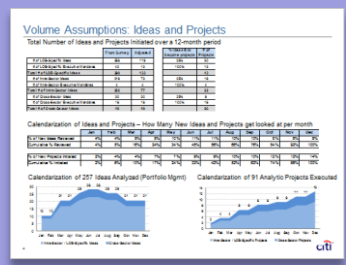


Templates for estimating and amortizing build and support costs for hardware, software, etc. services/resources

**PROCEDURES WIKI**

Online intranet/extranet communication and collaboration resources to support platform operations, BAU and break-fix operations

# Operating Framework on Data Lake Architecture

**Governance:**
Establish rules, policies and standards to protect, exploit and maximize the value of information in the organization

**Metadata:**
Define a business metadata strategy that is key in harmonizing information across disparate data sources and for consistent use of information by business users.

**Enterprise Models:**
Data need analysis, authoritative sources, standard data structures

**Data Integration:**
Evaluate data integration needs and make decisions around consistent use of EII, EAI,ETL

**Master Data Management:**
Provide a gold copy of reference data to the enterprise.

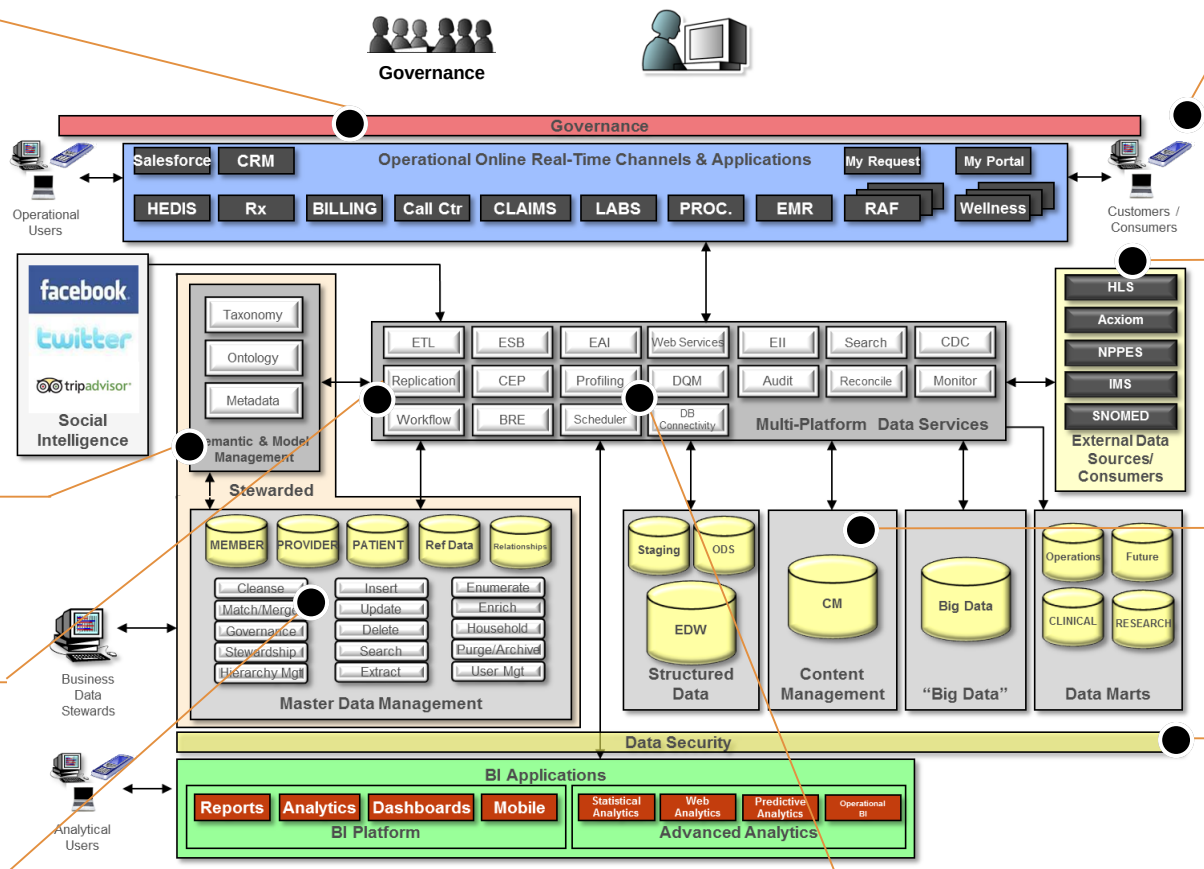


**Tools and Technologies:**
Standardizing tools and technologies based on best of breed tools.

**Data Quality:**
Implement on-going processes to measure and improve the timeliness, accuracy and consistency of the data.

**Presentation:**
Strategy to allow users to access information in a user-friendly manner.

**Data Access:**
Provide standard ways of sharing data with applications, business intelligence tools and downstream applications.

**Enterprise Content Mgmt:**
Provide a platform for delivery, storage and search for structured and un-structured data

**Security:**
Provide access to data based on roles using common technologies for access management and security management across different layers.