

---

# MITS6005

## Big Data

---

***Copyright © 2015 - 2019, Victorian Institute of Technology.***

*The contents contained in this document may not be reproduced in any form or by any means, without the written permission of VIT, other than for the purpose for which it has been supplied. VIT and its logo are trademarks of Victorian Institute of Technology.*

---

# Session 10 & 11

## Spark Deployments

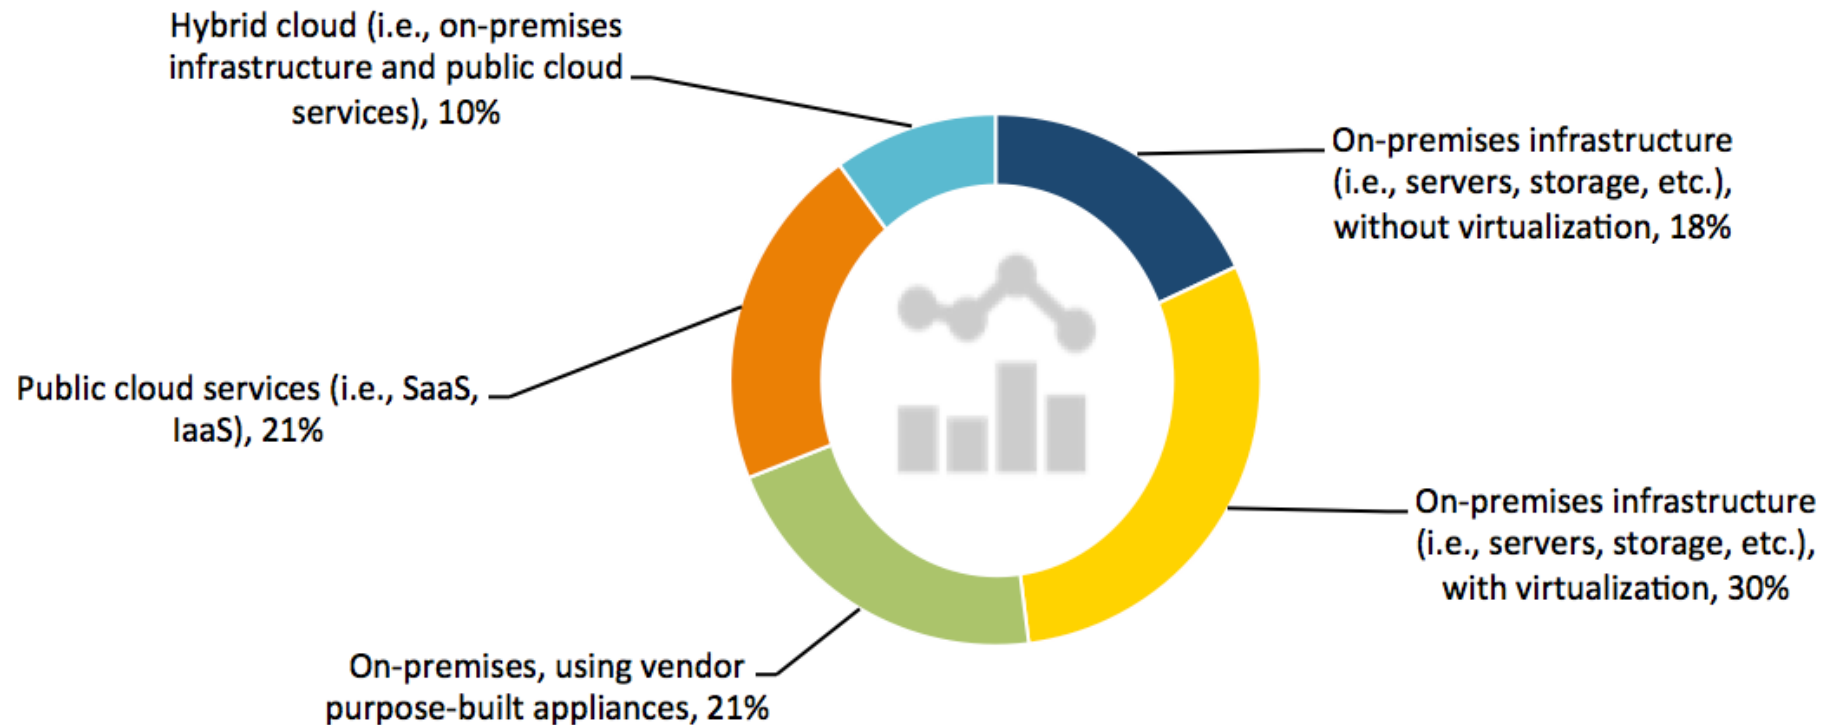
---

***Copyright © 2015 - 2019, Victorian Institute of Technology.***

*The contents contained in this document may not be reproduced in any form or by any means, without the written permission of VIT, other than for the purpose for which it has been supplied. VIT and its logo are trademarks of Victorian Institute of Technology.*

# Big Data Deployment Options


**In terms of net-new BI/analytics deployments, which of the following best describes the primary deployment strategy your organization will likely use going forward? (Percent of respondents, N=370)**



Source: Enterprise Strategy Group (ESG) Survey, 2015

# How Spark runs on clusters

- Spark applications run as independent sets of processes on a cluster, coordinated by the SparkContext object in your main program (called the *driver program*).
- To run on a cluster, the SparkContext can connect to several types of *cluster managers*, which allocate resources across applications.
- Once connected, Spark acquires *executors* on nodes in the cluster, which are processes that run computations and store data for your application. Next, it sends your application code (defined by JAR or Python files passed to SparkContext) to the executors.
- Finally, SparkContext sends *tasks* to the executors to run.



Source: <https://spark.apache.org/docs/latest/cluster-overview.html>

# Spark Cluster Management Types

- [Standalone](#) – a simple cluster manager included with Spark that makes it easy to set up a cluster.
- [Apache Mesos](#) – a general cluster manager that can also run Hadoop MapReduce and service applications.
- [Hadoop YARN](#) – the resource manager in Hadoop 2.
- [Kubernetes](#) – an open-source system for automating deployment, scaling, and management of containerized applications.

## Most Common Spark Deployment Environments (Cluster Managers)



**48%**

Standalone mode



**40%**

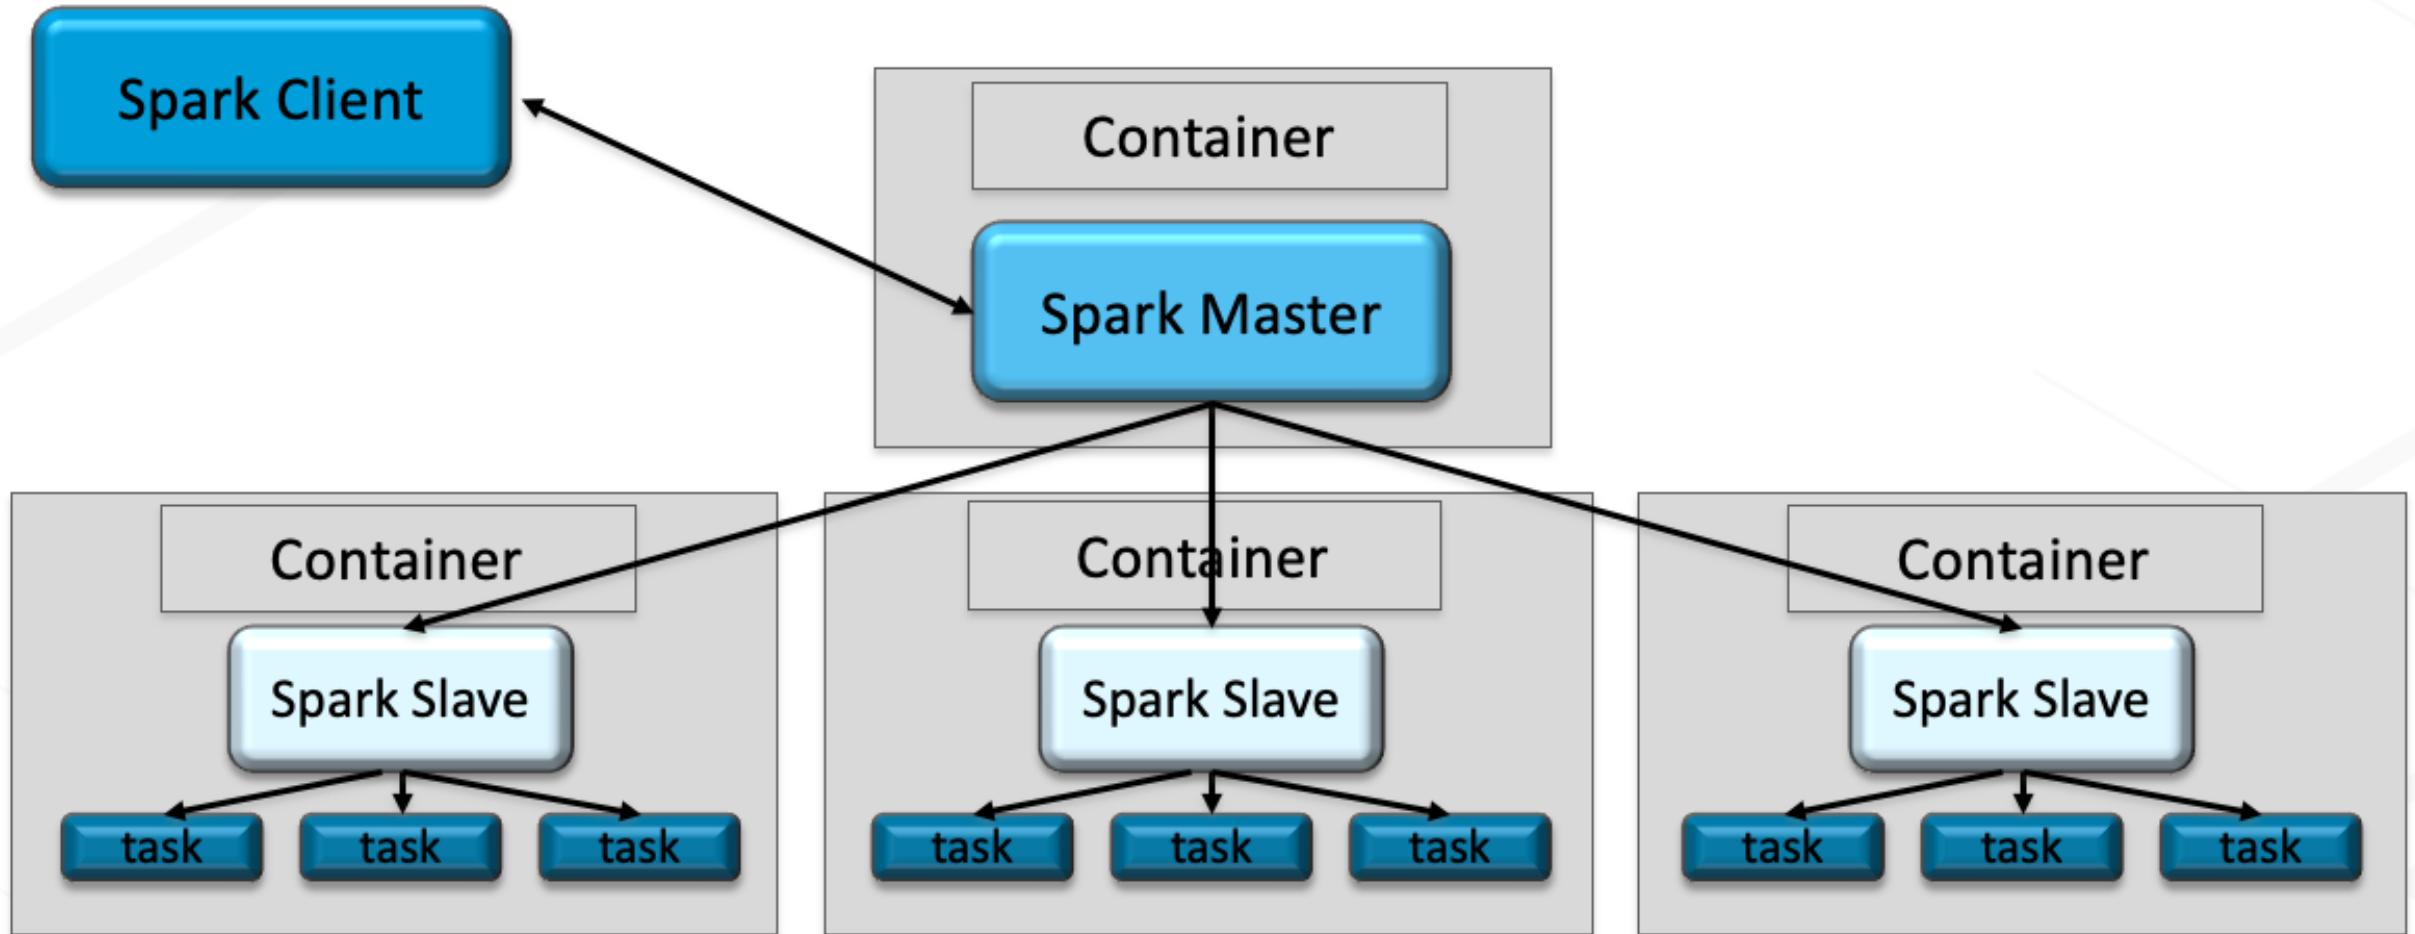
YARN



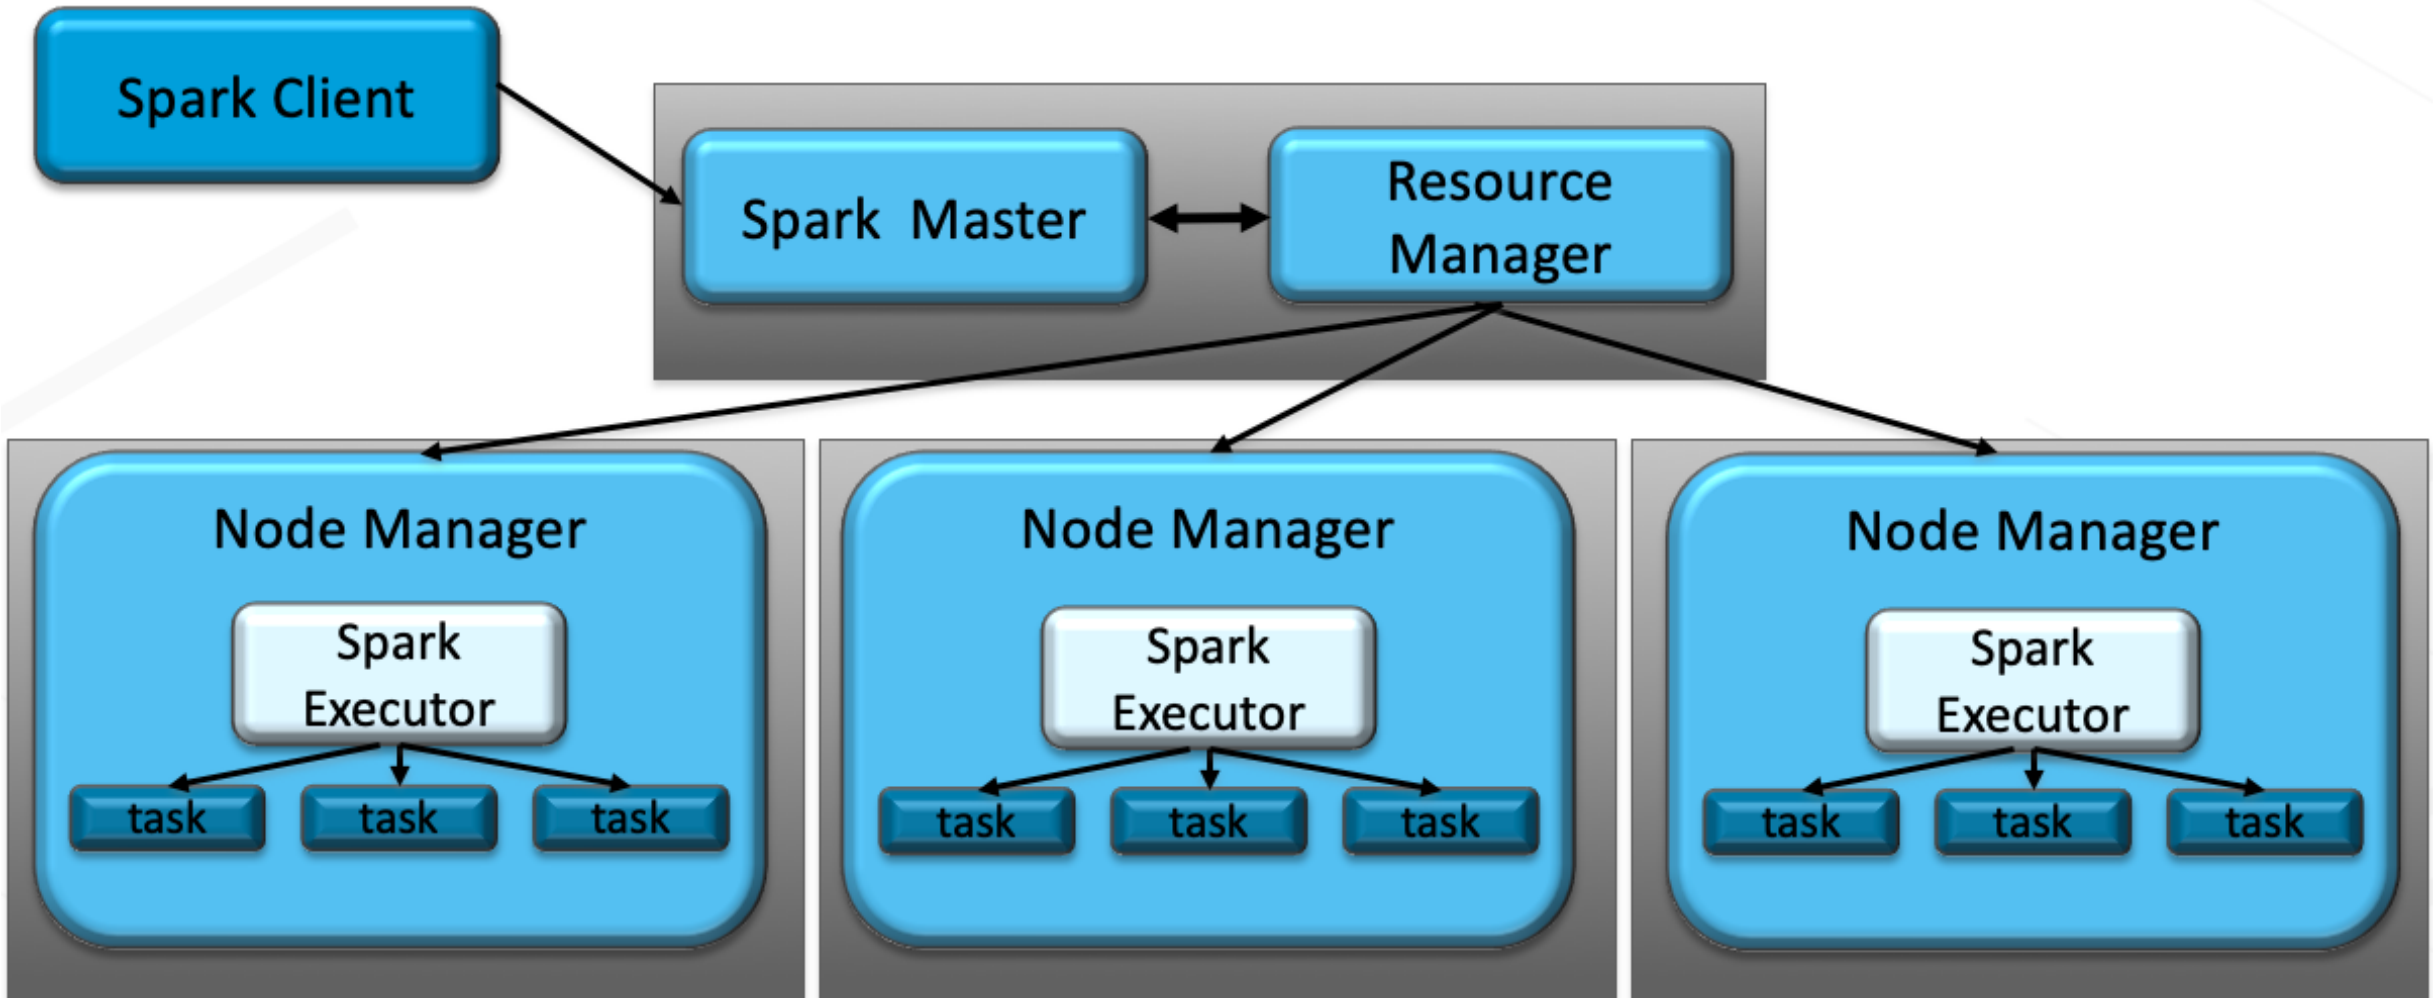
**11%**

Mesos

# Spark Cluster – Standalone



# Spark Cluster – Hadoop YARN






# What is Real-Time Data Analysis?

- Spot risk as it happens across multiple data sets
- Tap into data that is always on (sensor data, machine logs, web logs, etc.)
- React to changing business conditions in real time
- Spot opportunities before it is too late

# Visualizing Different Analytics

## Data Marts



- Data warehouse
- High response time
- Low latency
- Purpose-built
- No room for innovation

## Data Lake



- Batch data
- High latency
- Iterative
- Exploratory
- Try a few different ways

## Real-



- Real-time data
- Low latency
- High throughput
- More unknowns
- No room for error

# Real-Time Analysis Use Cases

applications	sensors	web	mobile phones
intrusion detection	malfunction detection	site analytics	network metrics analysis
fraud detection	dynamic process optimisation	recommendations	location based ads
log processing	supply chain planning	sentiment analysis	...

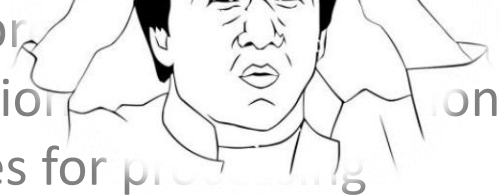
# What is Spark Streaming?

- Framework for large scale stream processing
  - Scales to 100s of nodes
  - Can achieve second scale latencies
  - Integrates with Spark's batch and interactive processing
  - Provides a simple batch-like API for implementing complex algorithm
  - Can absorb live data streams from Kafka, Flume, ZeroMQ, etc.

# Case study: XYZ, Inc.

- Any company who wants to process live streaming data has this problem
- Twice the effort to implement any new function
- Twice the number of bugs to solve
- Twice the headache

- Two processing stacks



Custom-built distributed stream processing engine

- 100% custom code
- Complex machine learning on millions of data points
- Requires many dozens of nodes for processing




Hadoop based for batch analysis

- General purpose daily and monthly reports
- Similar computation as the streaming system

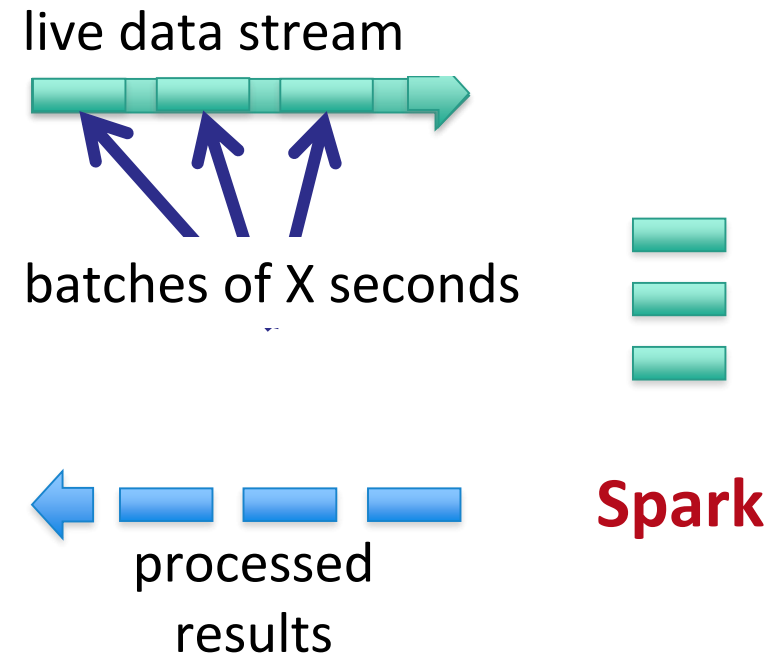
# Stateful Stream Processing

- Traditional streaming systems have a event-driven **record-at-a-time** processing model
  - Each node has mutable state
  - For each record, update state & send new records
- State is lost if node dies!
- Making stateful stream processing be fault-tolerant is challenging



Run a streaming computation as a **series of very small, deterministic batch jobs**


- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches





Run a streaming computation as a **series of very small, deterministic batch jobs**

- Batch sizes as low as  $\frac{1}{2}$  second, latency  $\sim 1$  second
- Potential for combining batch processing and streaming processing in the same system





# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
```

**DStream:** a sequence of RDD representing a stream of data

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream



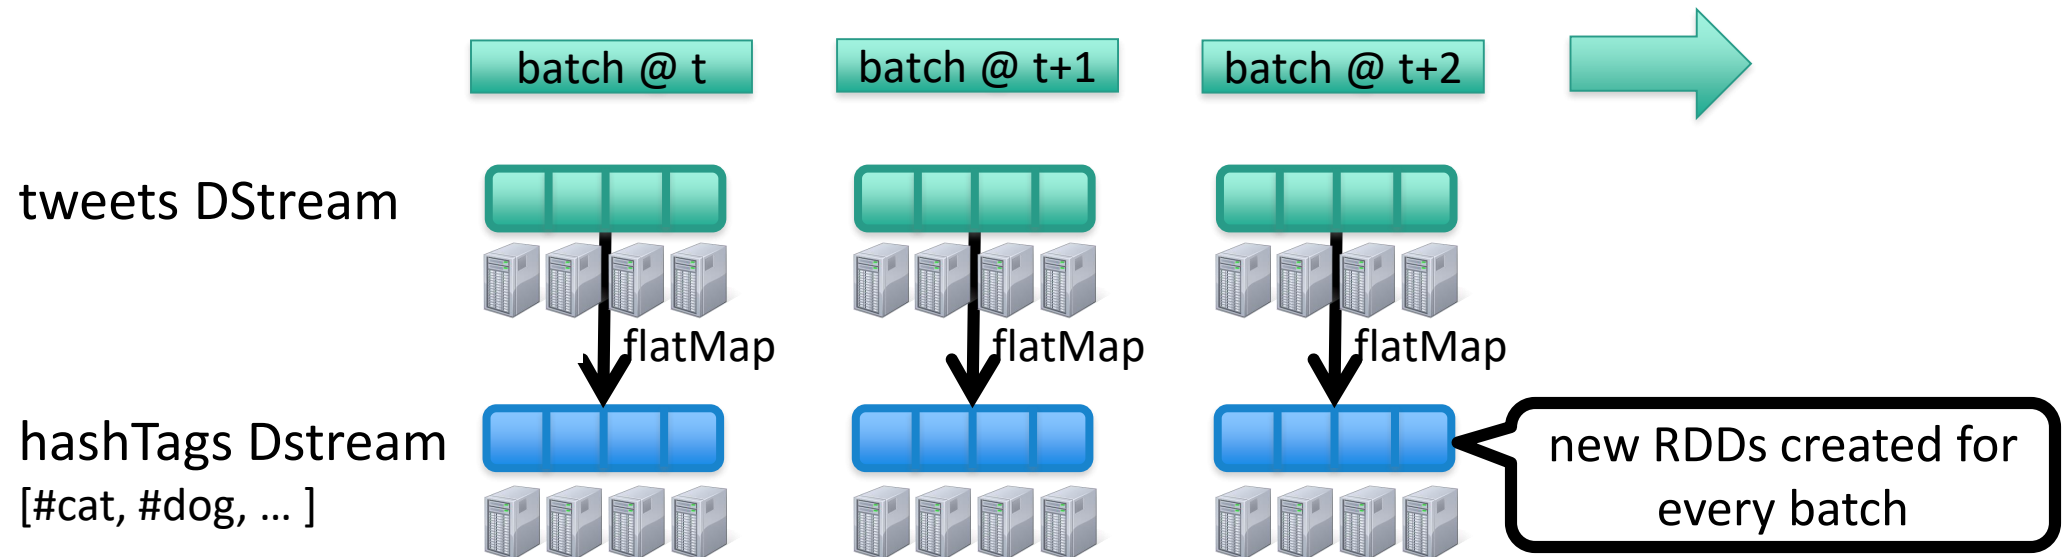
stored in memory as an RDD  
(immutable, distributed)

# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)  
val hashTags = tweets.flatMap(status => getTags(status))
```

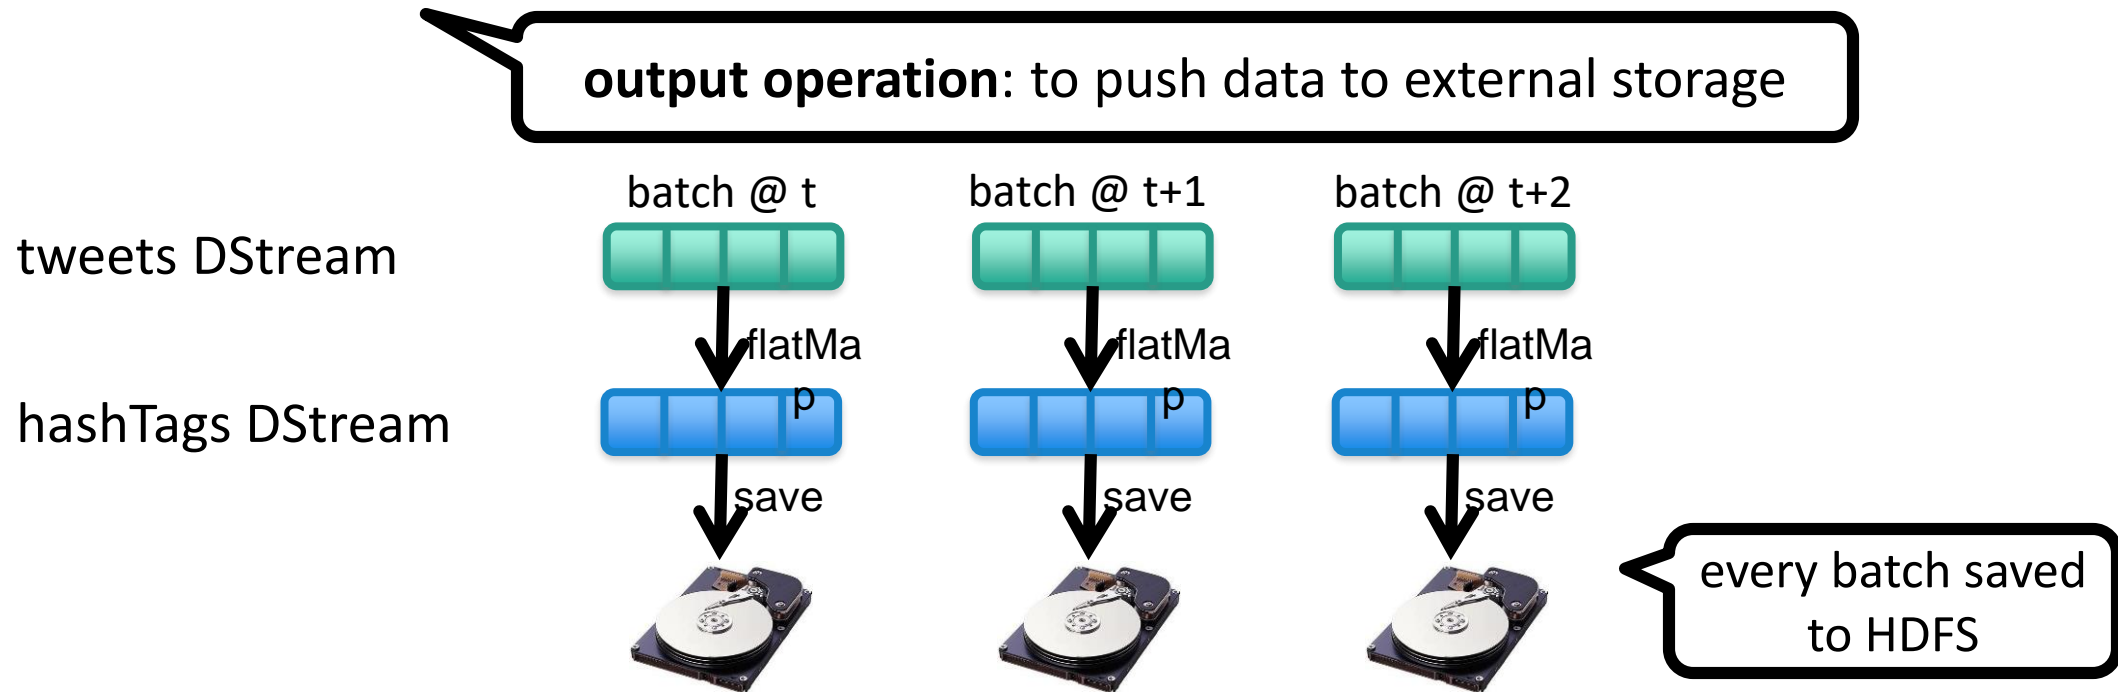
new DStream

**transformation:** modify data in one Dstream to create another DStream




# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)  
val hashTags = tweets.flatMap (status => getTags(status))  
hashTags.saveAsHadoopFiles("hdfs://...")
```



# Fault-tolerance


- RDDs remember the sequence of operations that created it from the original fault-tolerant input data
- Batches of input data are replicated in memory of multiple worker nodes, therefore fault-tolerant
- Data lost due to worker failure, can be recomputed from input data



- **DStream** – sequence of RDDs representing a stream of data
  - Twitter, HDFS, Kafka, Flume, ZeroMQ, Akka Actor, TCP sockets
- **Transformations** – modify data from on DStream to another
  - Standard RDD operations – map, countByValue, reduce, join, ...
  - Stateful operations – window, countByValueAndWindow, ...
- **Output Operations** – send data to external entity
  - saveAsHadoopFiles – saves to HDFS
  - foreach – do anything with each batch of results

## Example 2 – Count the hashtags

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)  
val hashTags = tweets.flatMap (status => getTags(status))  
val tagCounts = hashTags.countByValue()
```



## Example 3 – Count the hashtags over last 10 mins

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)  
val hashTags = tweets.flatMap (status => getTags(status))  
val tagCounts = hashTags.window(Minutes(10), Seconds(1)).countByValue()
```

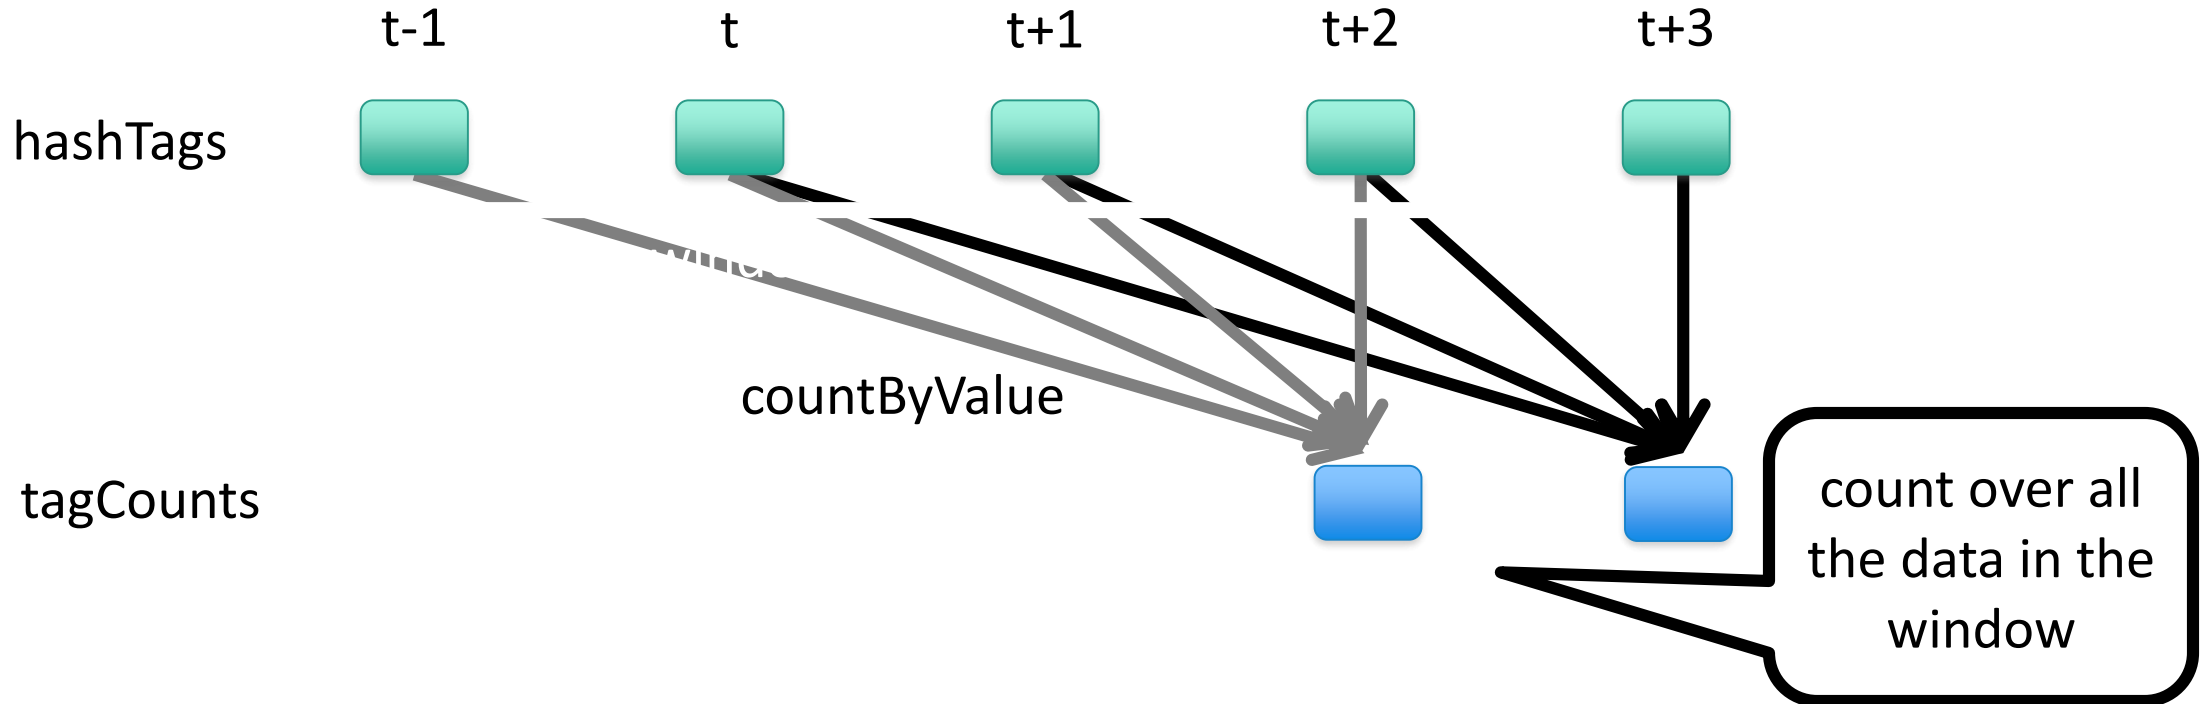
sliding window  
operation

window length

sliding interval

## Example 3 – Counting the hashtags over last 10 mins

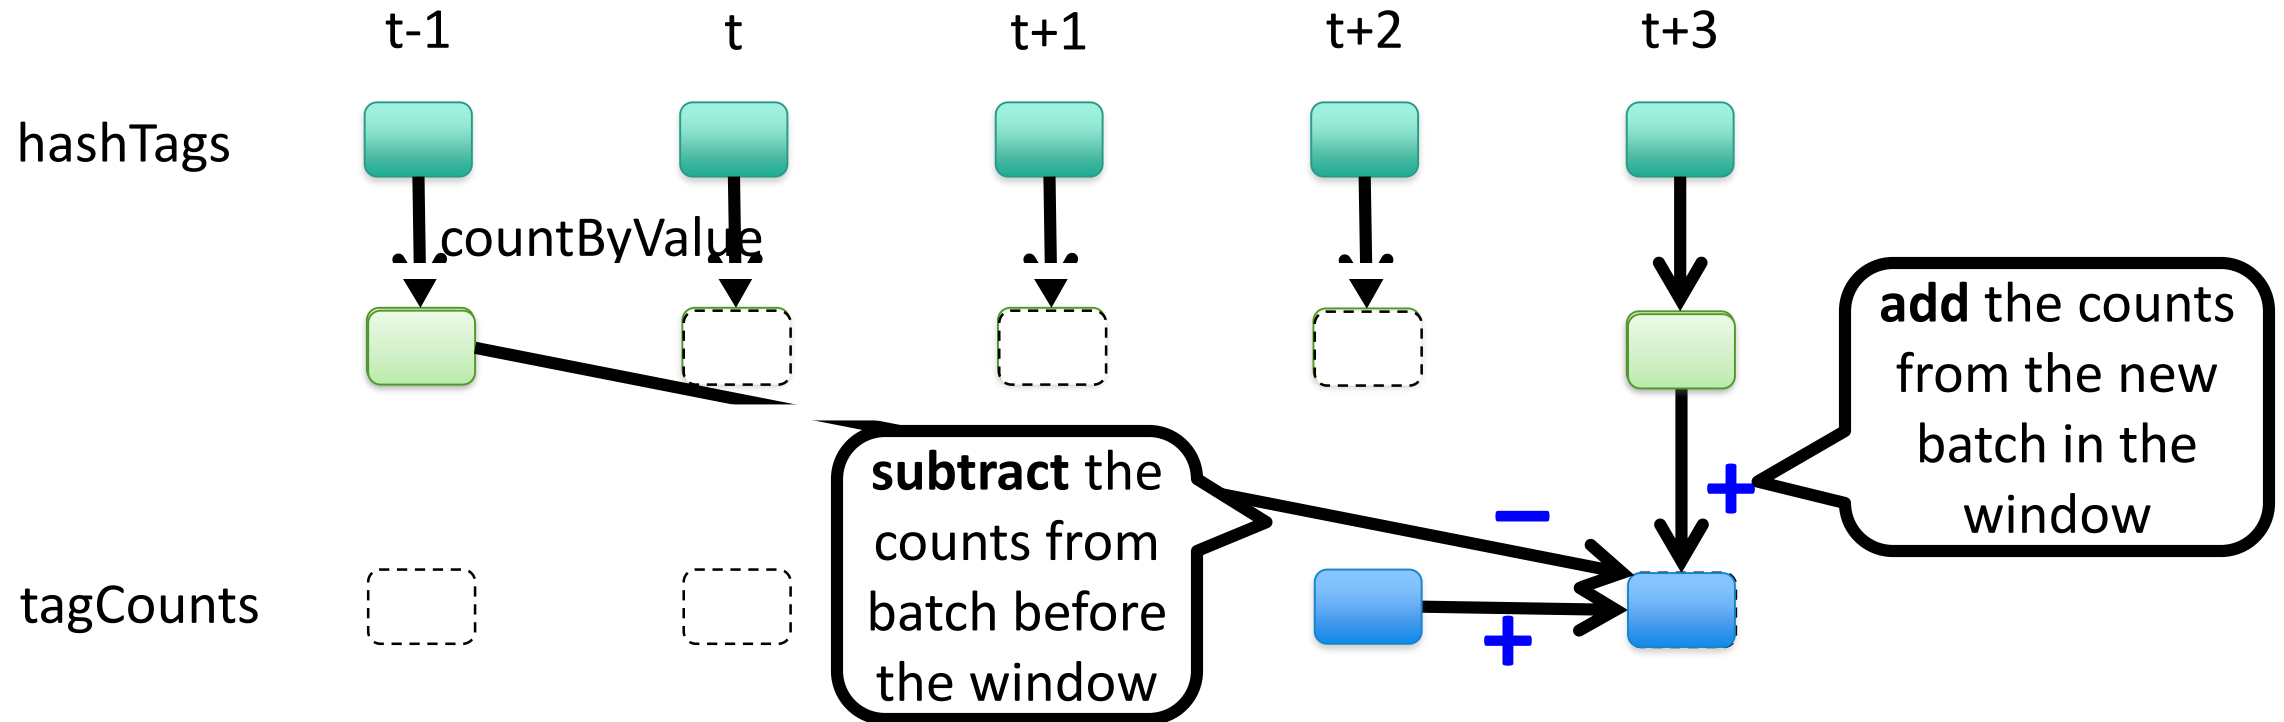
```
val tagCounts = hashTags.window(Minutes(10), Seconds(1)).countByValue()
```





# Smart window-based countByValue

```
val tagCounts = hashtags.countByValueAndWindow(Minutes(10), Seconds(1))
```




- Technique to incrementally compute count generalizes to many reduce operations
  - Need a function to “inverse reduce” (“subtract” for counting)
- Could have implemented counting as:

```
hashTags.reduceByKeyAndWindow(_ + _, _ - _, Minutes(1), ...)
```

# Fault-tolerant Stateful Processing

All intermediate data are RDDs, hence can be recomputed if lost



- State data not lost even if a worker node dies
  - Does not change the value of your result
- *Exactly once* semantics to all transformations
  - No double counting!

- Maintaining arbitrary state, track sessions
  - Maintain per-user mood as state, and update it with his/her tweets

```
tweets.updateStateByKey(tweet => updateMood(tweet))
```

- Do arbitrary Spark RDD computation within DStream
  - Join incoming tweets with a spam file to filter out bad tweets

```
tweets.transform(tweetsRDD => {  
    tweetsRDD.join(spamHDFSFile).filter(...)  
})
```

# An example - Sentiment Analysis with PySpark

First step in any Apache programming is to create a SparkContext.

```
import findspark
findspark.init()
import pyspark as ps
import warnings
from pyspark.sql import SQLContext

try:
    # create SparkContext on all CPUs available: in my case I have 4 CPUs on my laptop
    sc = ps.SparkContext('local[4]')
    sqlContext = SQLContext(sc)
    print("Just created a SparkContext")
except ValueError:
    warnings.warn("SparkContext already exists in this scope")
```

For full details on sentiment analysis: <https://towardsdatascience.com/sentiment-analysis-with-pyspark-bc8e83f80c35>

